

Transformers and GNNs

CPSC483: Deep Learning on Graph-Structured Data

Rex Ying

Readings

- Readings are updated on the website (syllabus page)
- **Lecture 7 readings:**
 - [Graph Attention Networks](#)
 - [Multi-hop Attention Graph Neural Networks](#)
- **Lecture 8 readings:**
 - [Attention is All You Need](#)
 - [Graph Structure of Neural Networks](#)

Outline of Today's Lecture

1. Self-Attention and Transformers

2. Transformers Applications

3. Graph Transformers and Sparse Transformers

Sequence Learning

- Inputs from different domains can be seen as the general **sequence** of **tokens**

Domain

Sequence

Token

Structure

NLP

Sentence: [SOS, “graph”,
“neural”, “networks”, “are”,
“powerful”, EOS]


Word: “graph”
Phrase: [“graph”,
“neural”, “networks”]

Sequential correlations

CV

Image:






Pixel
Patch: 

Spatial correlations

Graph

Graph:



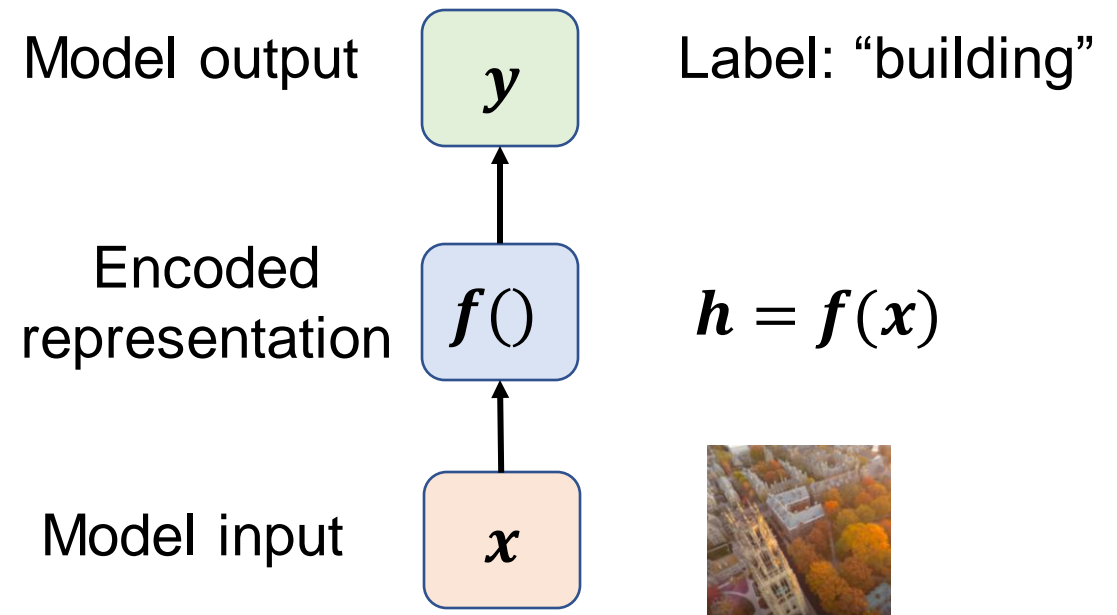
Node: 
Subgraph: 

Adjacency

Standard Supervised Learning Setting

- **One (token) to One (token)**

- Input is a single token (e.g., an image), and the output is its attribute (e.g., label) or another token.
- $h = f(x)$, $f()$ is the model to learn.



Sequence Learning — Application

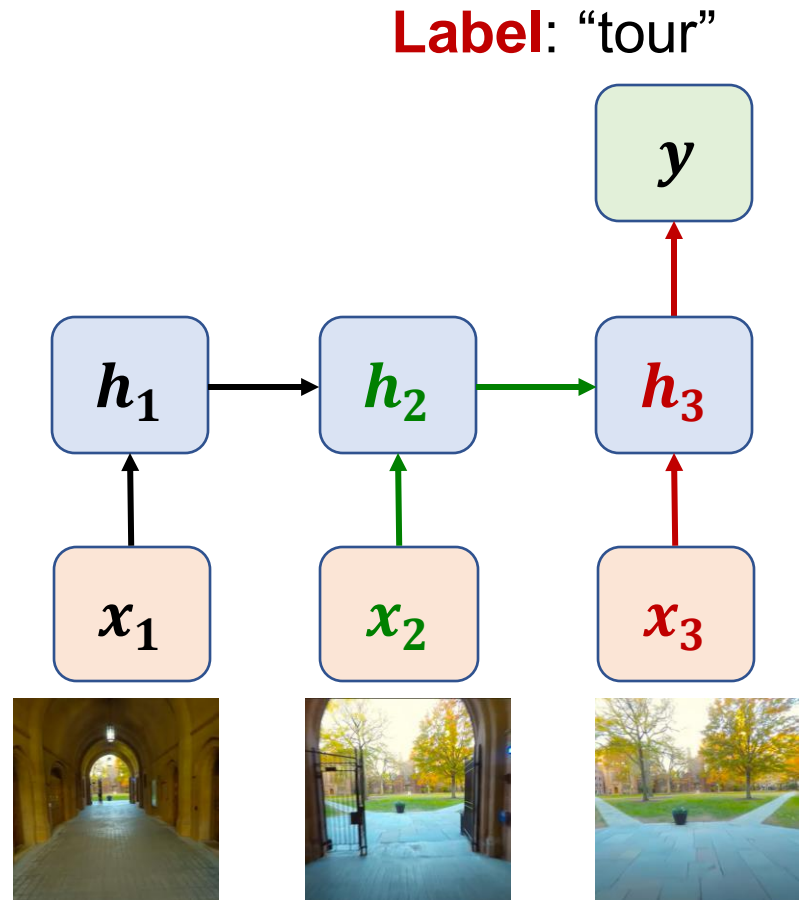
- **Many (tokens) to One**

- Input is a sequence of tokens (e.g. a video with frames), and the output is its attribute (e.g. label) or another token.
- $h_1 = f(x_1)$
- To generate h_2 , we would like to incorporate both x_2 and the preceding frame x_1 $h_2 = f(x_2, h_1)$, and $f()$ is still the same model (shared parameters)

- $h_i = f(x_i, h_{i-1})$

Current token

Previous token




Sequence Learning — Application

- We can process a sequence of tokens $X = [x_1, x_2, \dots, x_n]$ by applying a recurrence formula at every time step
- **Recurrent neural networks**

$$h_i = f_W(x_i, h_{i-1})$$

↓ new state ↓ current input ↘ old state

- For example, $h_i = \sigma(W_x x_i + W_h h_{i-1} + b_h)$,
and $y_i = \sigma(W_y h_i + b_y)$



A folded diagram of RNNs

Sequence Learning — Application

- **Many (tokens) to Many**

- The sequence is first encoded into a hidden representation, then gradually decoded by the decoder.

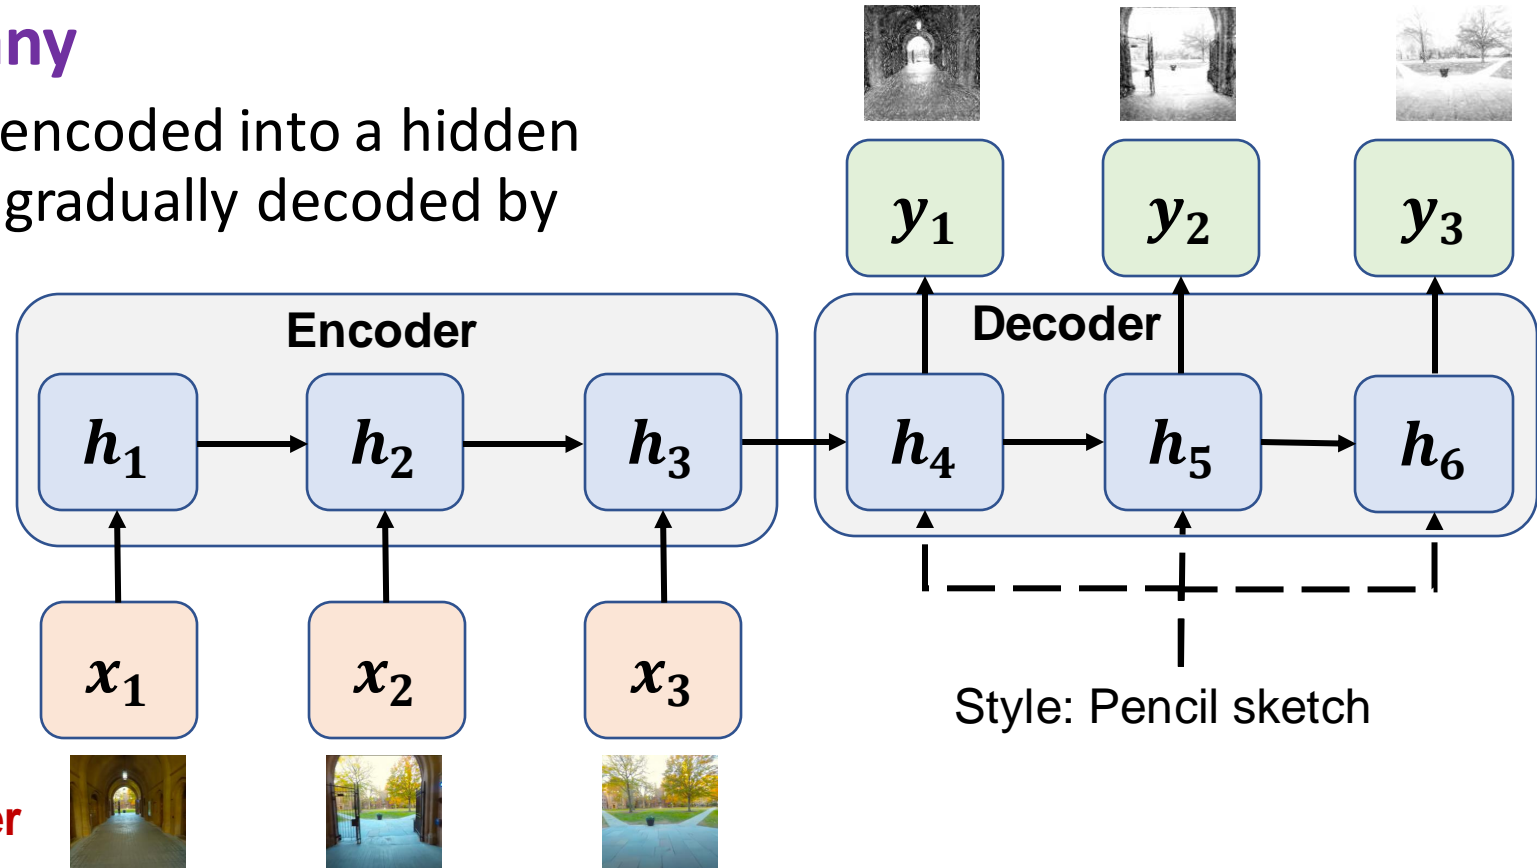


Diagram of video **style transfer**

Sequence Learning — Application

- **One (token) to One (token)**

- **Many to One**


- Protein to property
- Sentence to sentiment
- Song to name

- **One to Many**

- Image to caption

- **Many to many**

- English to Chinese
- Time series: history to future
- Graph autoregression



Speech classification
(many-to-one)

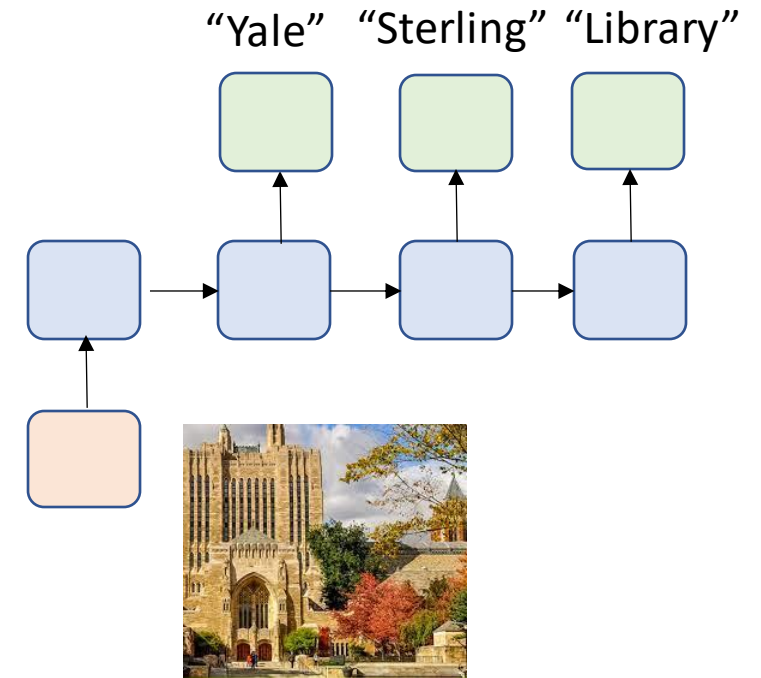



Image Captioning
(one-to-many)

Sequence Learning




$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$


Problems of RNNs

- Sequential computation prevents parallelization
- Capacity of handling long sequences
- Mainly focusing on modeling recurrence
 - does not capture other correlations (hierarchical, long-range, polysemy....) well

Sequence Learning



$$h_i = f_W(x_i, h_{i-1}), y_i = f_Y(h_i)$$



$$y_i = \text{self-att}([x_i + p_i]_i)$$

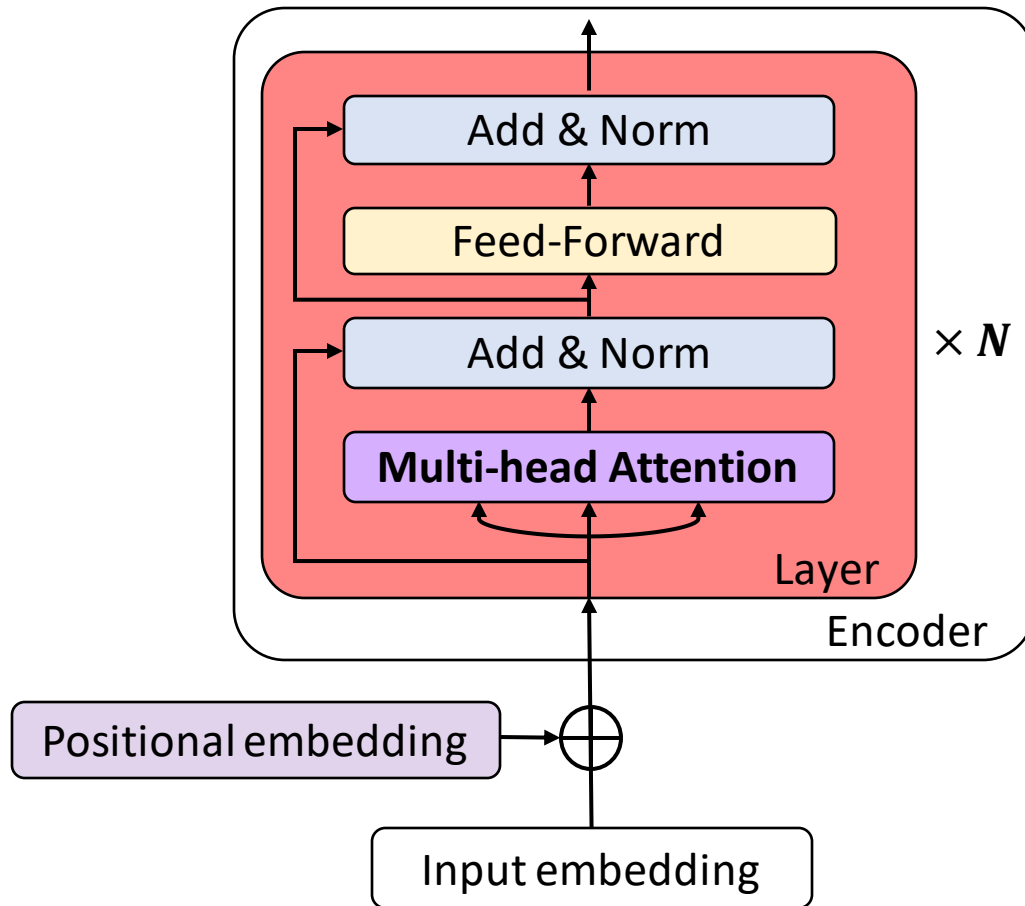
Problems of RNNs

1. parallelization — — — — — ➤
2. long sequences — — — — — ➤
3. only recurrence — — — — — ➤

Solutions by Transformers


1. **Parallel input:** Input All tokens at the same time
2. **Self-Attention:** Enable attention in long-range
3. **Positional Embeddings p_i :** Model all possible correlations

Transformers — Overview



- **Original paper:** Attention is all you need [Vaswani et al., 2019].
- **Key component:** multi-head self-attention
- **Transformer layer:** Multi-Head Self-Attention, layer normalization, skip connection, position-wise FFN
- **Model usage:** Pretraining then fine-tune, with Encoder only/ Encoder-Encoder/ Decoder only (discussing later)

Transformers — Overview



- **Design choices** of transformers:
(for those interested in transformer architectures)

Absolute/relative position, equivariant embedding (for graph)

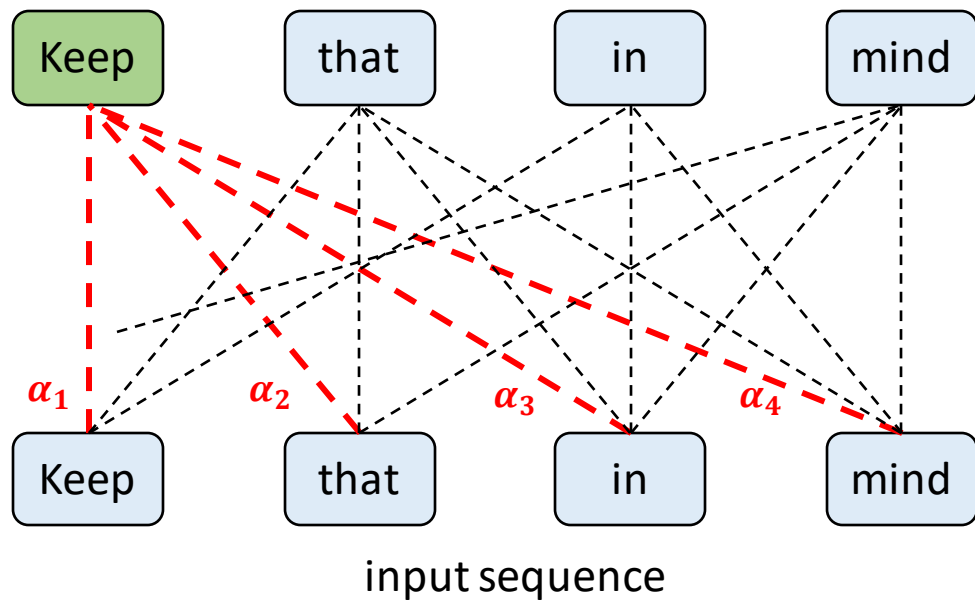
Sparse attention, low-rank attention, attention with prior, memory compression, query prototyping...

Placement, substitutes, normalization-free

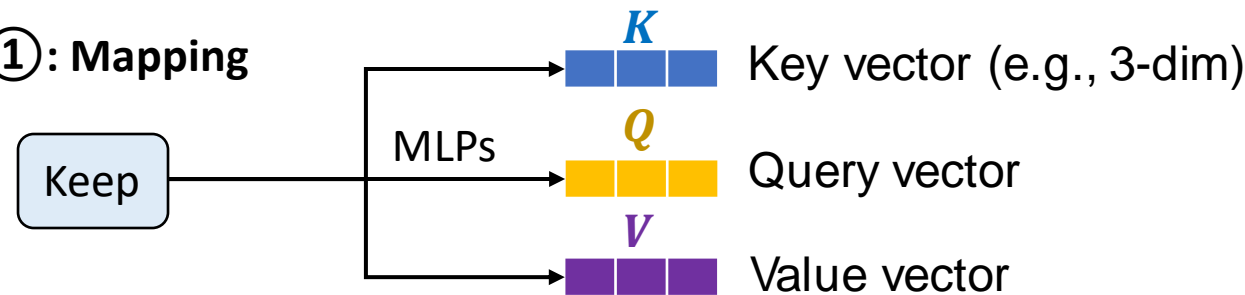
Cross-block connections, recurrence/hierarchy, other architecture

Transformers — Self-Attention (1/5)

Example:



Step ①: Mapping



Step ②: Attention



$$\alpha_1, \alpha_2, \alpha_3, \alpha_4 = \text{Softmax} \left(\begin{matrix} Q \\ \text{Keep} \end{matrix} \times \begin{matrix} K & K & K & K \\ \text{Keep} & \text{that} & \text{in} & \text{mind} \end{matrix} \right)$$

Step ③: Update

$$\begin{matrix} V' \\ \text{Keep} \end{matrix} = \alpha_1 \times \begin{matrix} V \\ \text{Keep} \end{matrix} + \alpha_2 \times \begin{matrix} V \\ \text{that} \end{matrix} + \alpha_3 \times \begin{matrix} V \\ \text{in} \end{matrix} + \alpha_4 \times \begin{matrix} V \\ \text{mind} \end{matrix}$$

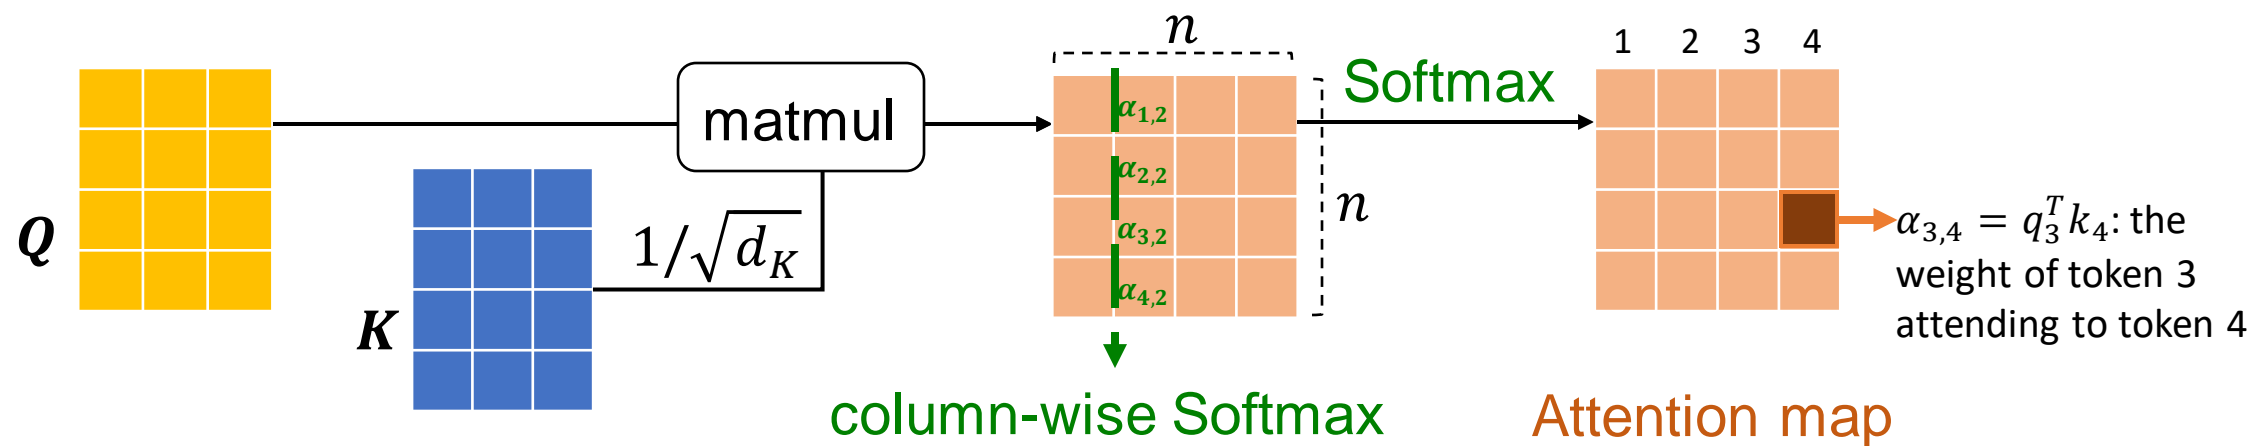
Transformers — Self-Attention (2/5)

- Formally, given an input sequence $X = [x_1, x_2, \dots, x_n] \in \mathbb{R}^{n \times d_X}$
- Step ①: Query $Q = XW_Q$, Key $K = XW_K$, Value $V = XW_V$
 - $W_K \in \mathbb{R}^{d_X \times d_K}$, and thus $K \in \mathbb{R}^{n \times d_K}$
 - We require $d_K = d_Q$, for simplicity, we set $d_K = d_Q = d_V := d$



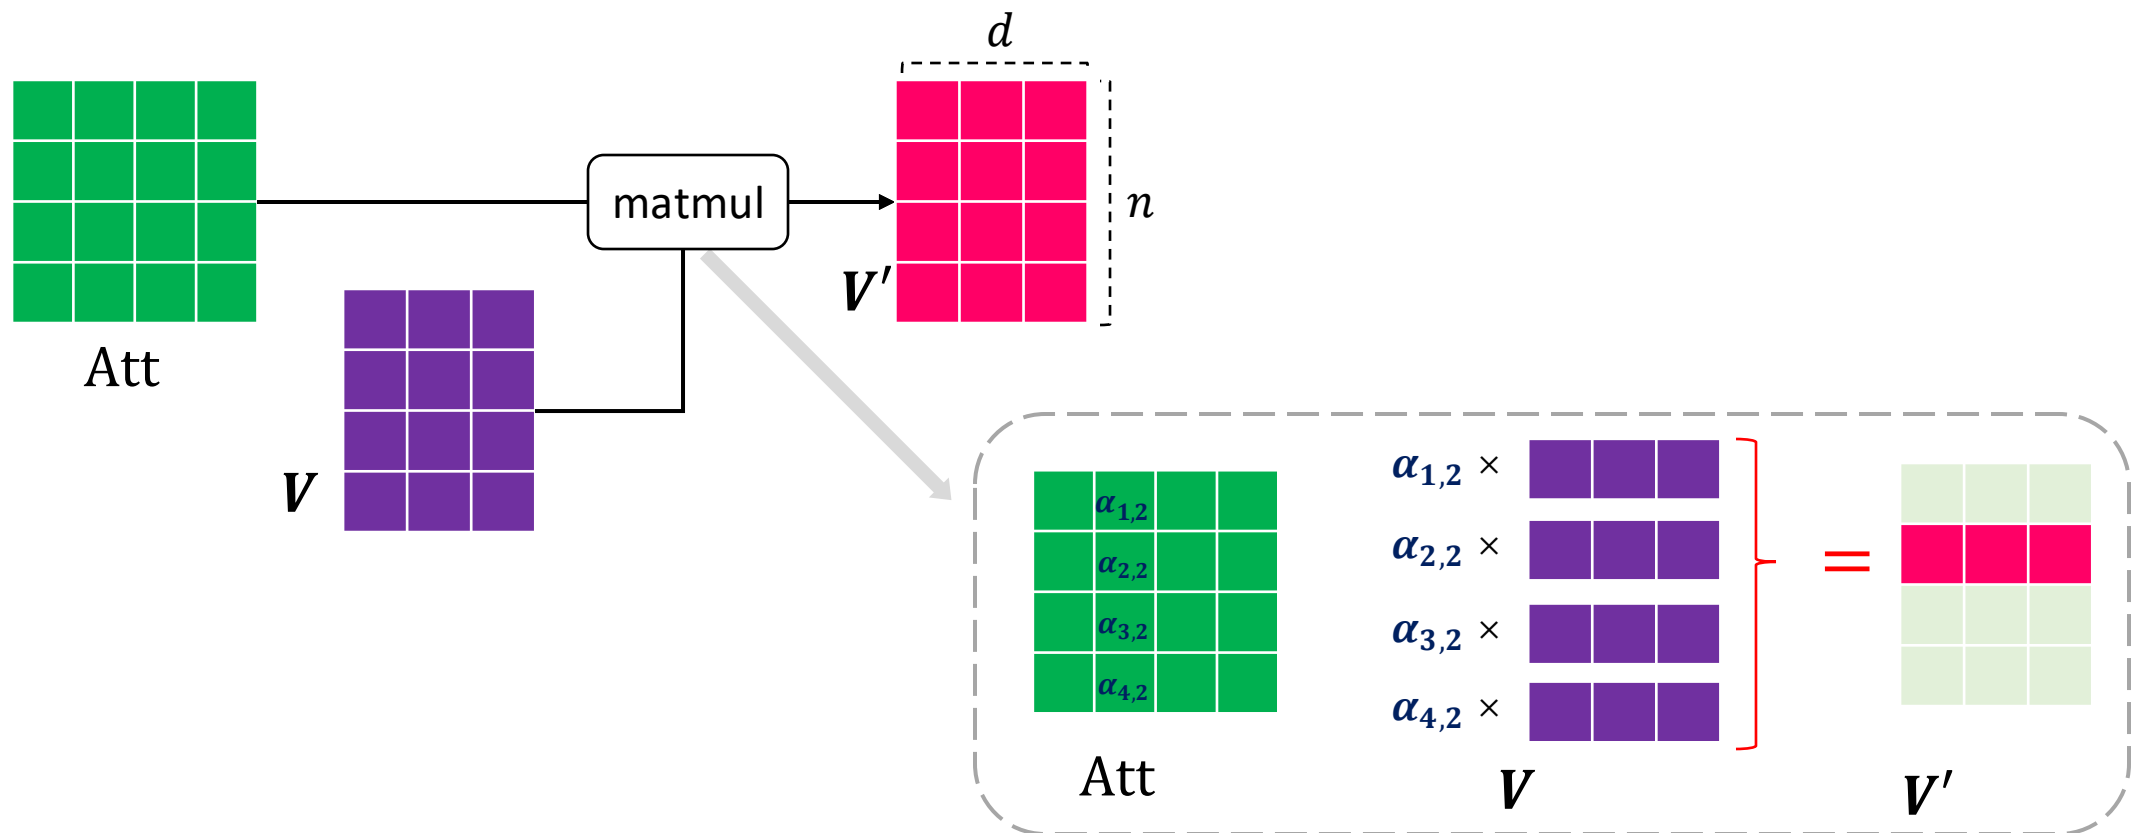
Transformers — Self-Attention (3/5)

- Step ② : Attention map $\text{Att} = \text{Softmax} \left(\frac{QK^T}{\sqrt{d}} \right) \in \mathbb{R}^{n \times n}$ (Softmax is col-wise)
 - The matrix multiplication QK^T performs dot-product for every possible pair of queries and keys, resulting in an attention map.
 - **Normalization factor** $1/\sqrt{d_K}$: performing dot-product over two vectors with variance σ^2 results in a scalar having d_K -times higher variance,
 - $q \sim N(0, \sigma^2), k \sim N(0, \sigma^2) \rightarrow \text{Var}(\sum_{i=1}^{d_K} q[i]k[i]) = \sigma^4 d_K$




Transformers — Self-Attention (4/5)

- Step ③: Updated value $V' = \text{Att } V \in \mathbb{R}^{n \times d}$ Matrix product



Transformers — Self-Attention (5/5)

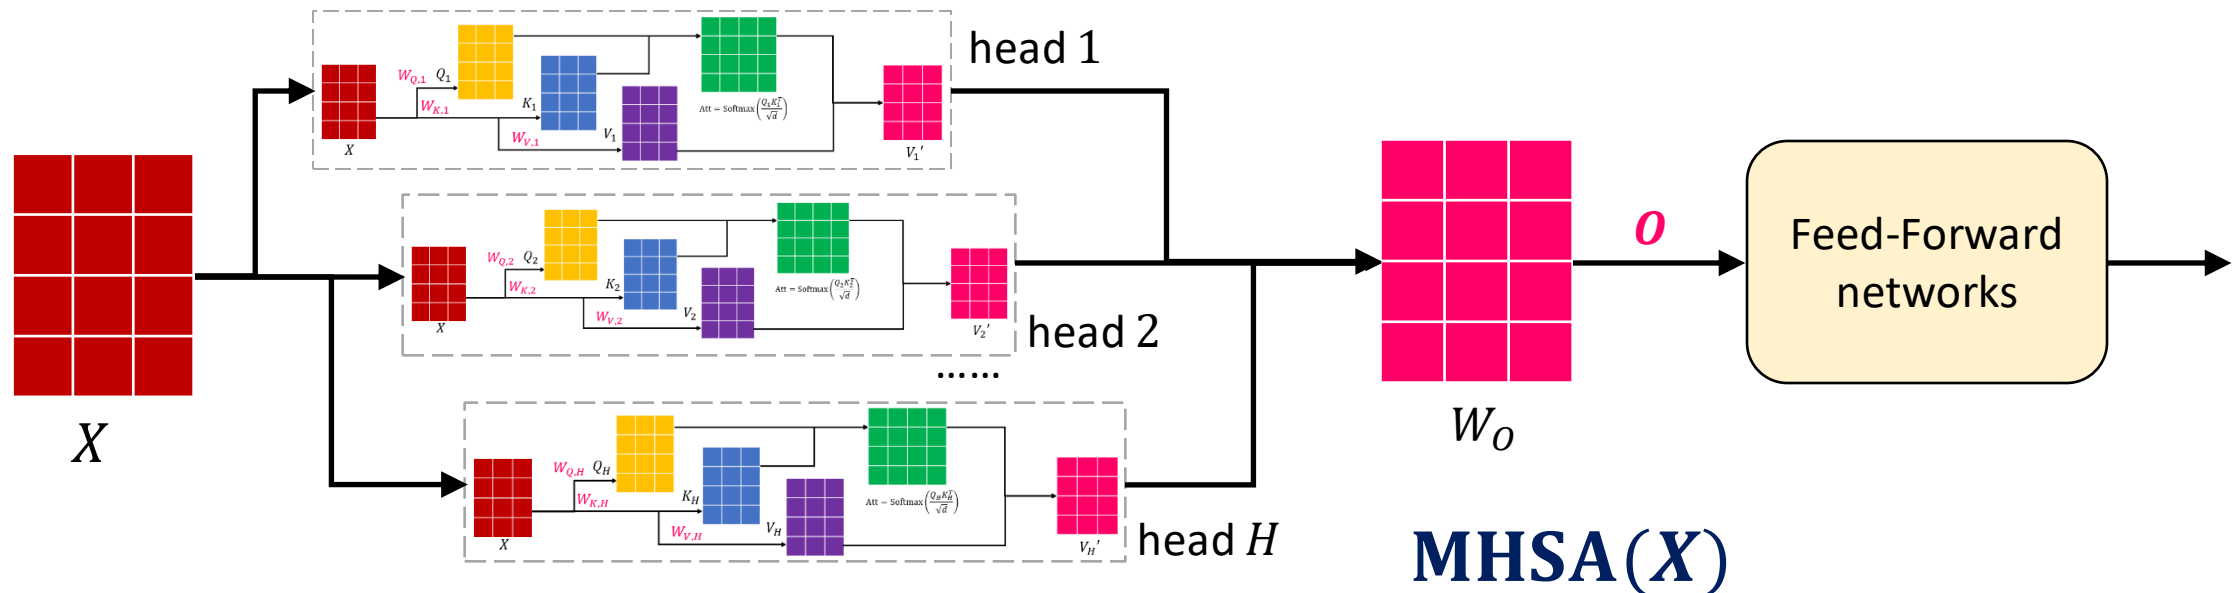
- Putting all together



The computation complexity is quadratic to number of tokens


Transformers — Multi-Head Self-Attention

- There are usually **multiple aspects** that a token can attend to.
- We extend the attention to multiple heads, with multiple (Q, K, V) triplets on the same features.
 - The output of multi-head self-attention $O = \text{Concat}([V_1', V_2', \dots, V_H'])W_O$
 - Learnable parameters in each attention layer: $W_{Q,i}, W_{K,i}, W_{V,i} \in R^{d_x \times d}$ for head i , $W_O \in R^{Hd \times d_o}$




Transformers — Layer (1)

- **MHSA**: multi-head self-attention
- Transformer layer: $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X))$
- **Residual connections** are added to
 - Enable smooth gradient flow in deep transformers
 - Keep the information of the original sequence.

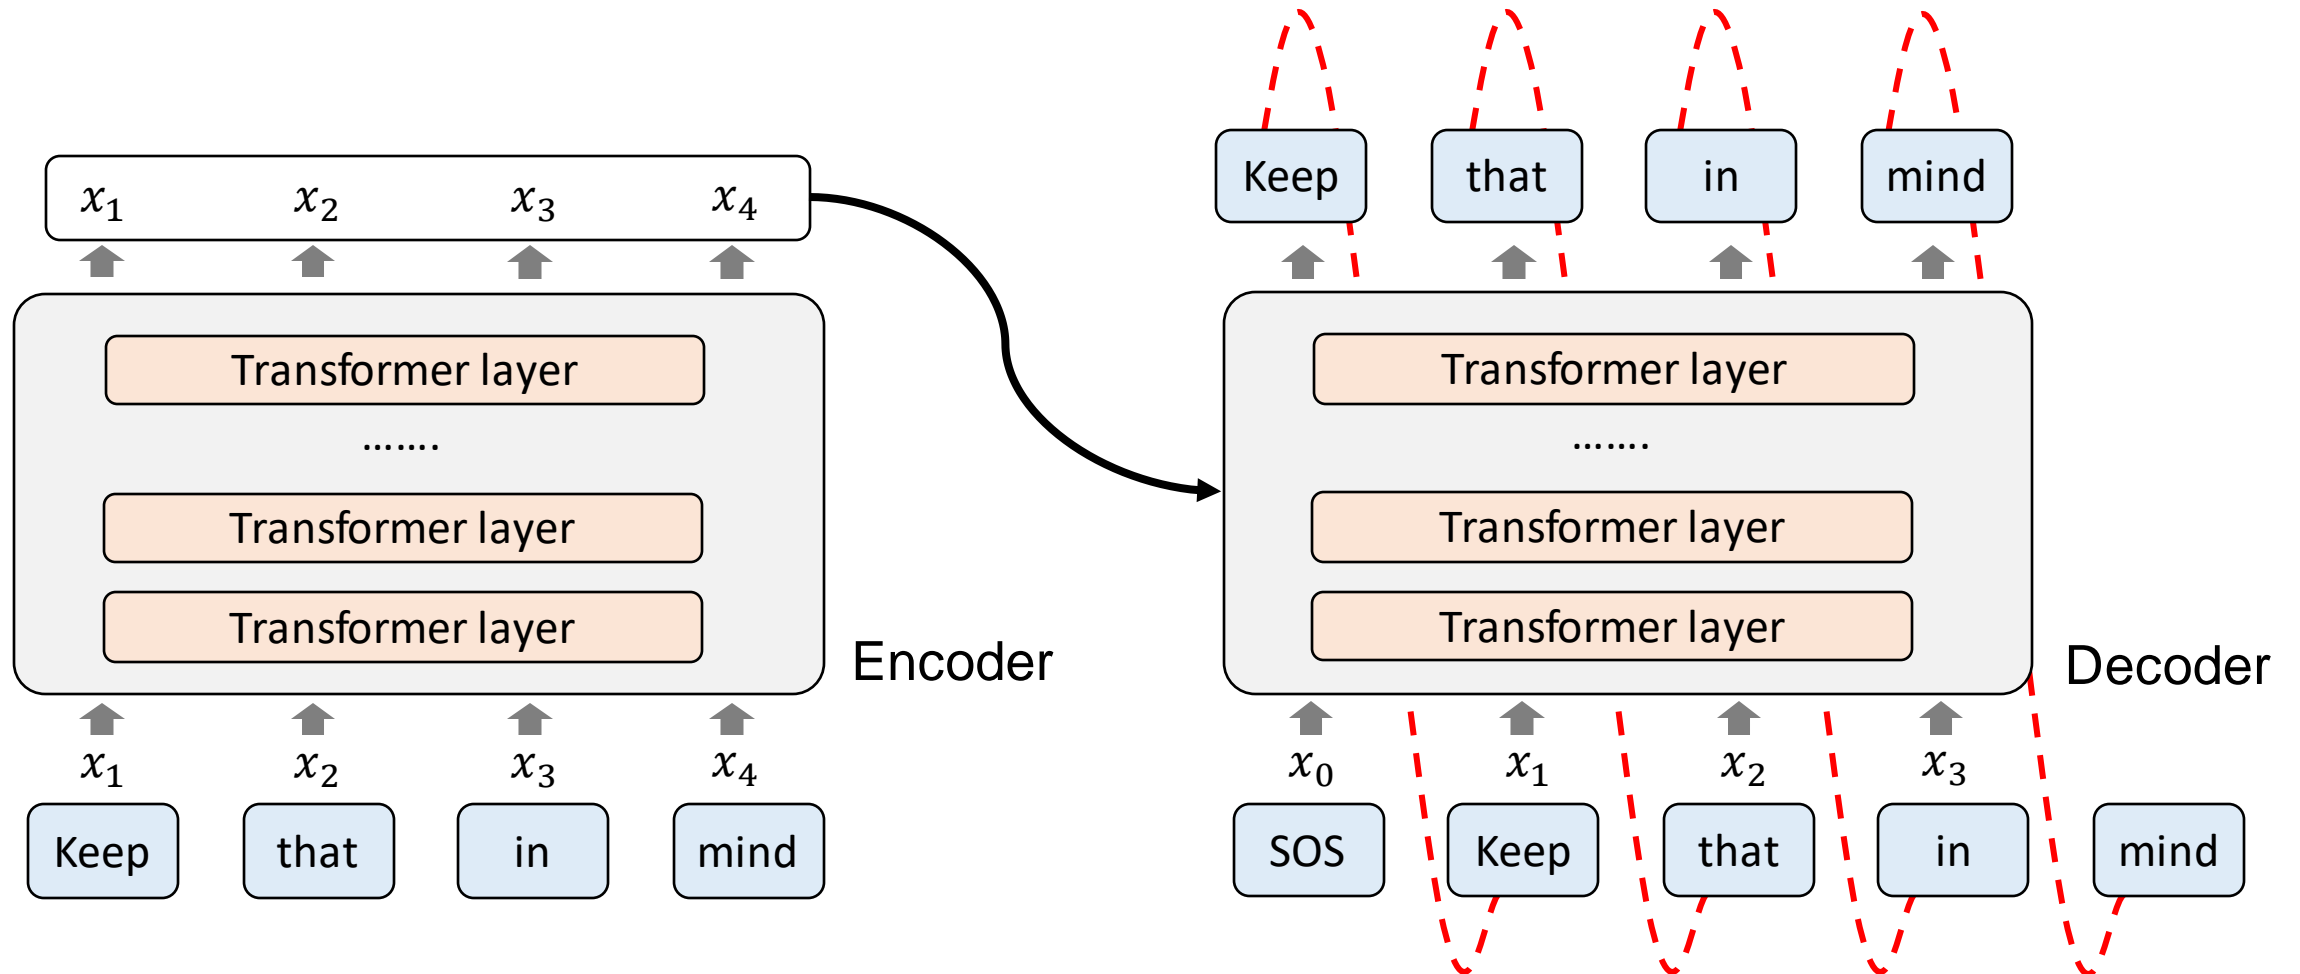


Transformers — Layer (2)

- Transformer layer: $X \rightarrow \text{LayerNorm}(X + \text{MHSA}(X)) \rightarrow \text{LayerNorm}(X + \text{FFN}(X))$
- Layer Normalization** is used to enable faster training with small regularization and keep features in similar magnitudes.
 - BatchNorm isn't applied because batch size is usually small in Transformers due to GPU memory constraints. Besides, BatchNorm has been shown to lead to worse performance in NLP.
- MLPs** are added for “post-processing”, and allow transformations on each sequence token.

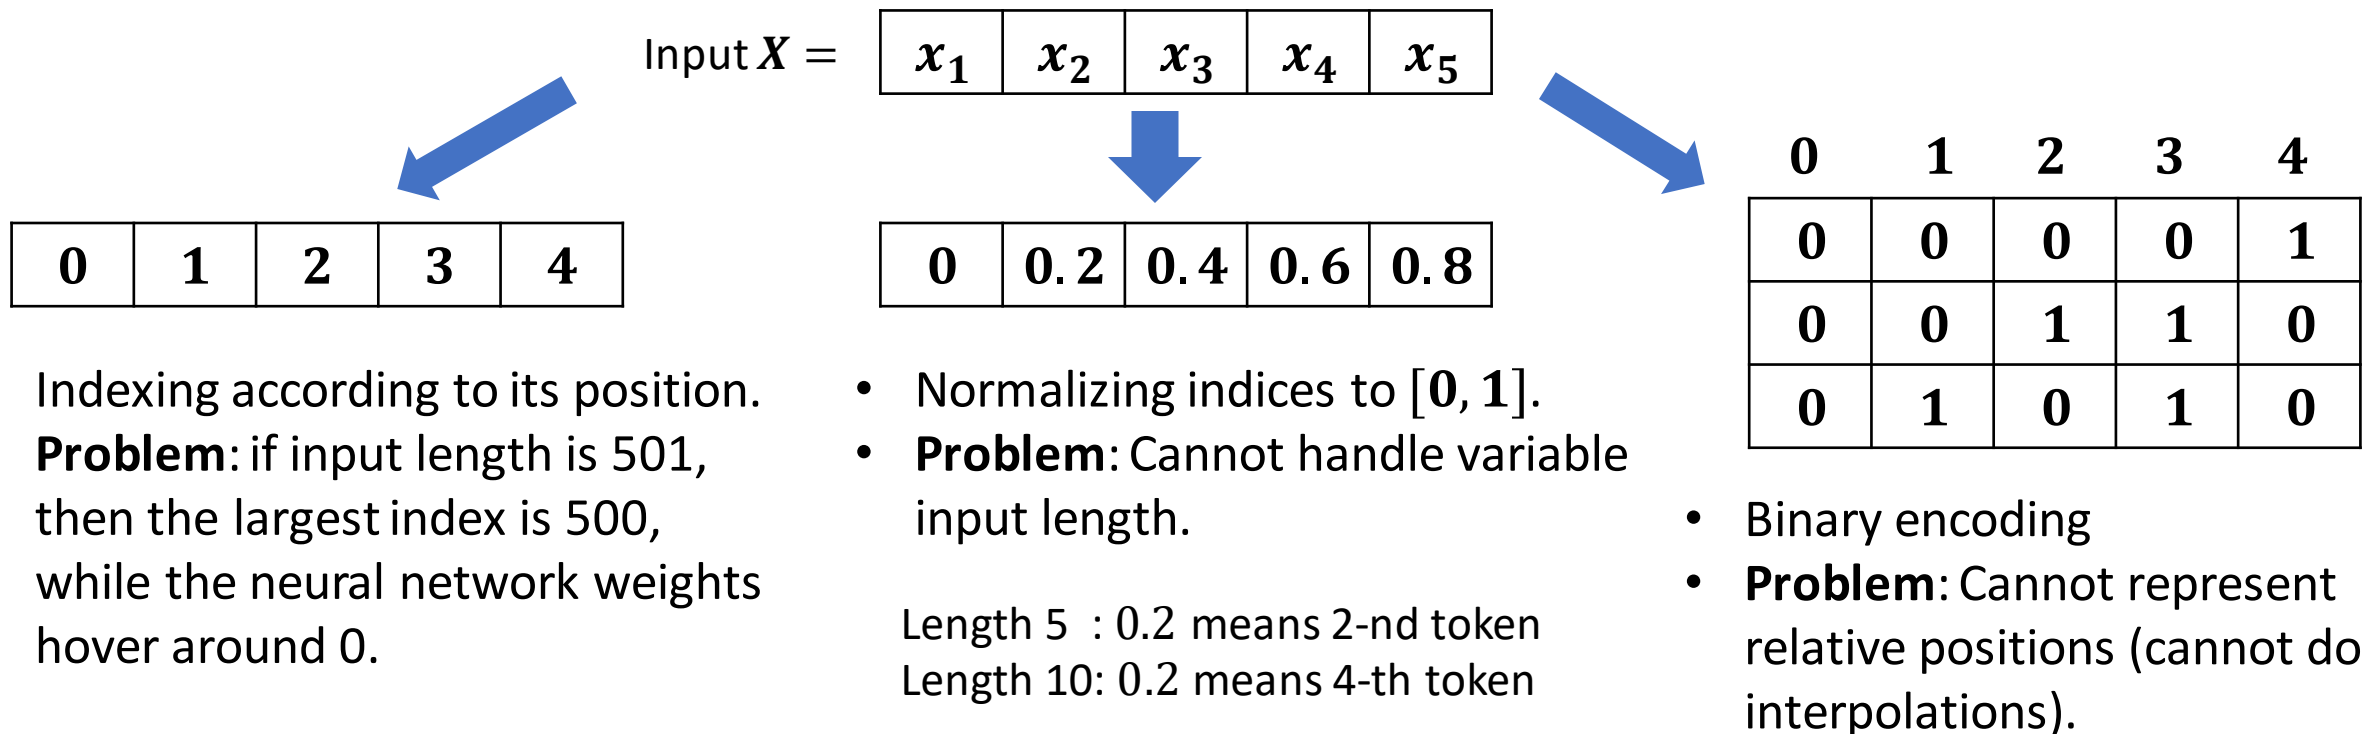


Transformers —Encoder / Decoder



Transformers —Positional Encoding (1)


Positional encoding is adopted to indicate the position of the token in the sequence.



We need discretization of something continuous!

Transformers —Positional Encoding (2)

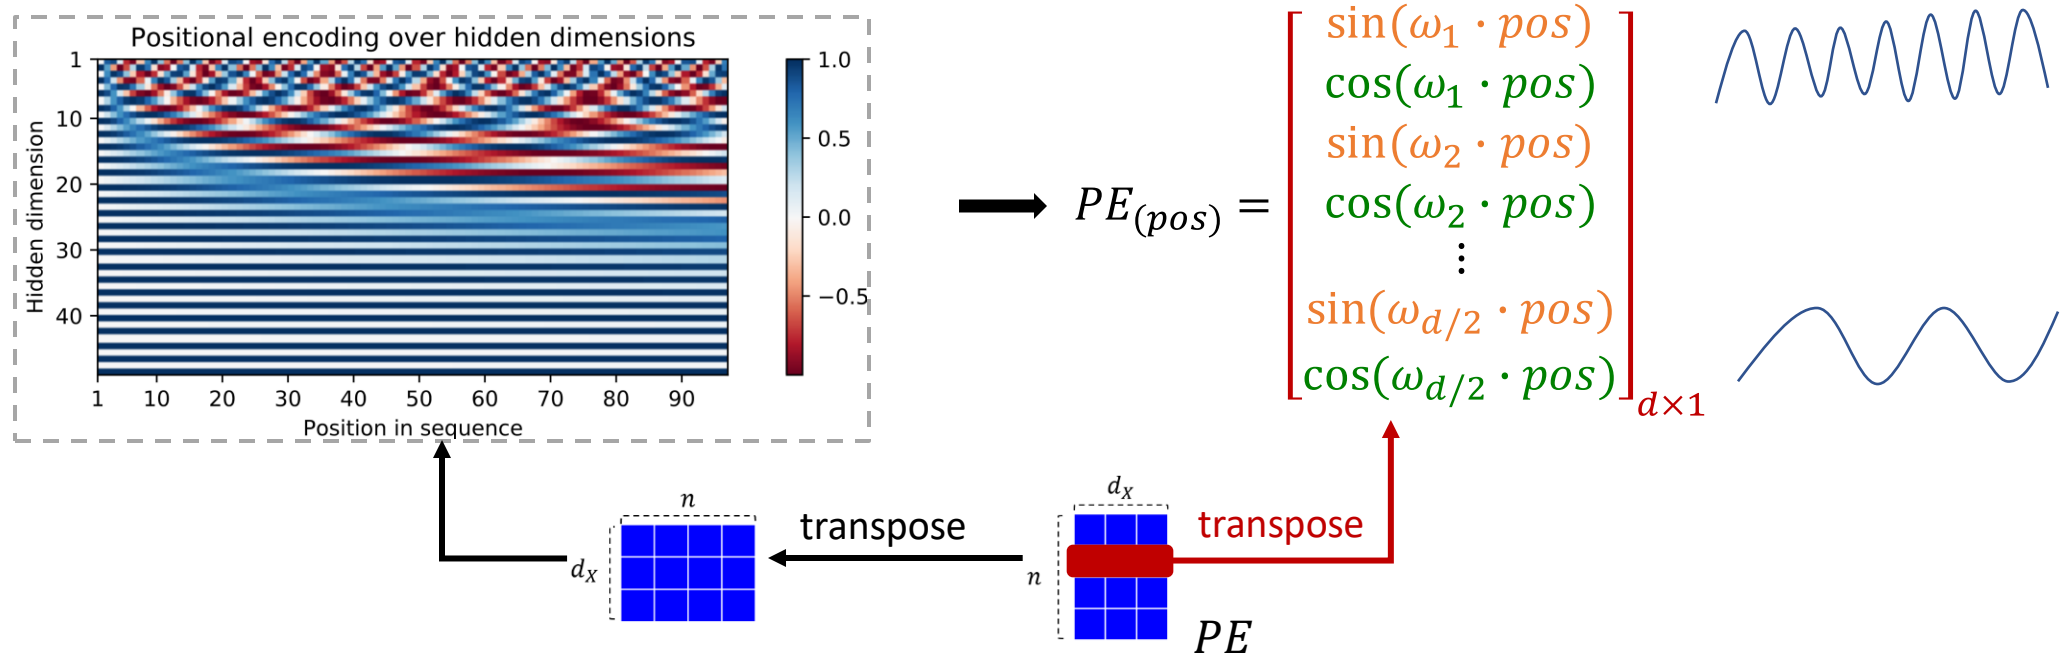
- **Learnable** 1-D/2-D embedding (commonly adopted in NLP/ CV):
 - We hope the relative positional relationship can be learnt through **back-propagation**.
- **Cosine encoding** (commonly adopted in NLP):
 - Recall the **input shape is of $\mathbb{R}^{n \times d_x}$** , we want the **PE has the same shape $\mathbb{R}^{n \times d_x}$**
 - For the **pos -th token (pos -th row)**,
 - Even columns: For **$2i$ -th dimension (column)**, we encode as $PE_{(pos,2i)} = \sin(pos/10000^{2i/d_x})$
 - Odd columns: For **$2i+1$ -th dimension (column)**, we encode as $PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_x})$



Transformers —Positional Encoding (3)

- **Cosine encoding**

- $PE_{(pos,2i)} = \sin(pos/10000^{2i/d_x})$, $PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_x})$.
- $\omega_i = 1/10000^{2i/d_x}$.
- Relative distance: $PE_{(pos+k)}$ can be easily represented as a linear function of $PE_{(pos)}$ (show it).




Summary: Transformer Architecture

- Multi-Head Self-Attention (**MHSA**(X))
 - For head i
 - $Q_i = XW_{Q_i}, K_i = XW_{K_i}, V_i = XW_{V_i}$
 - $\text{Att}_i = \text{Softmax}\left(\frac{Q_i K_i^T}{\sqrt{d}}\right) \in \mathbb{R}^{n \times n}$
 - $V_i' = \text{Att}_i V_i \in \mathbb{R}^{n \times d}$
 - Concatenating all heads: $O = \text{Concat}([V_1', V_2', \dots, V_H'])W_O$
- $X = \text{LayerNorm}(X + \text{MHSA}(X))$
- $X = \text{LayerNorm}(X + \text{FFN}(X))$

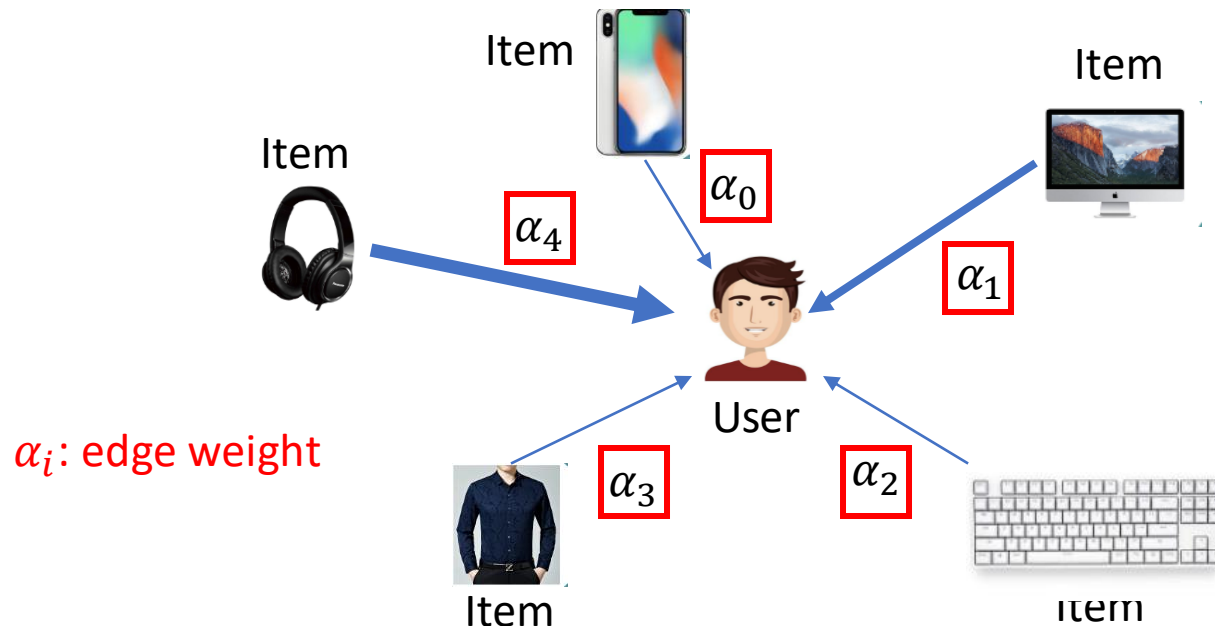
Recall: Graph Attention Mechanism (1)

- Let att be an **attention mechanism**
 - Attention coefficient e_{vu} is computed by a based on the messages of v, u :
$$e_{vu} = att(\mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}, \mathbf{W}^{(l)} \mathbf{h}_v^{(l-1)})$$
 - e_{vu} indicates the importance of u 's message to node v



Recall: Graph Attention Mechanism (2)

- Learnable weighting function can provide a good **interpretability**
 - Different edge weights indicates difference importance of the neighbor nodes
 - Take recommender system as an example



- User aggregates the information from items with different weights
 - High attention weights indicate that user prefers these corresponding items

Recall: Graph Attention Network (GAT)

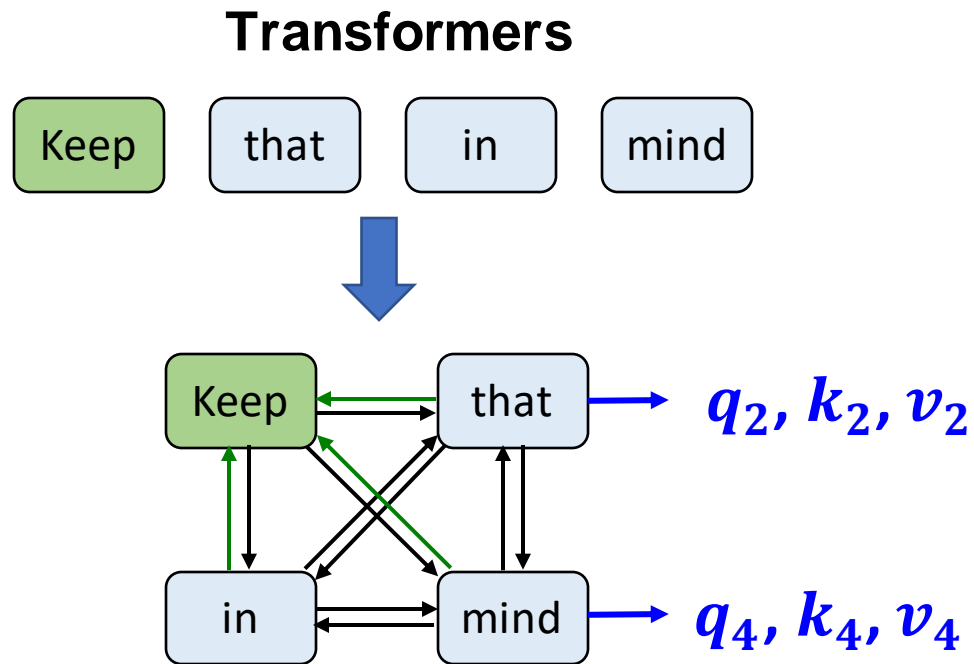
- A GAT layer (single head):
 - **Attention** computing: calculate the importance of neighbors
$$\alpha_{vu} = att(\mathbf{h}_v^{(l-1)}, \mathbf{h}_u^{(l-1)})$$
 - **Message** computing: transform information of neighbor node to a message
$$\mathbf{m}_u^{(l)} = \alpha_{vu} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}, u \in N_v$$
 - **Aggregate** message: aggregate messages from neighbor nodes
$$\mathbf{h}_v^{(l)} = \sigma \left(\sum_{u \in N_v} \mathbf{m}_u^{(l)} \right)$$

Learnable single-head or multi-head attention mechanism

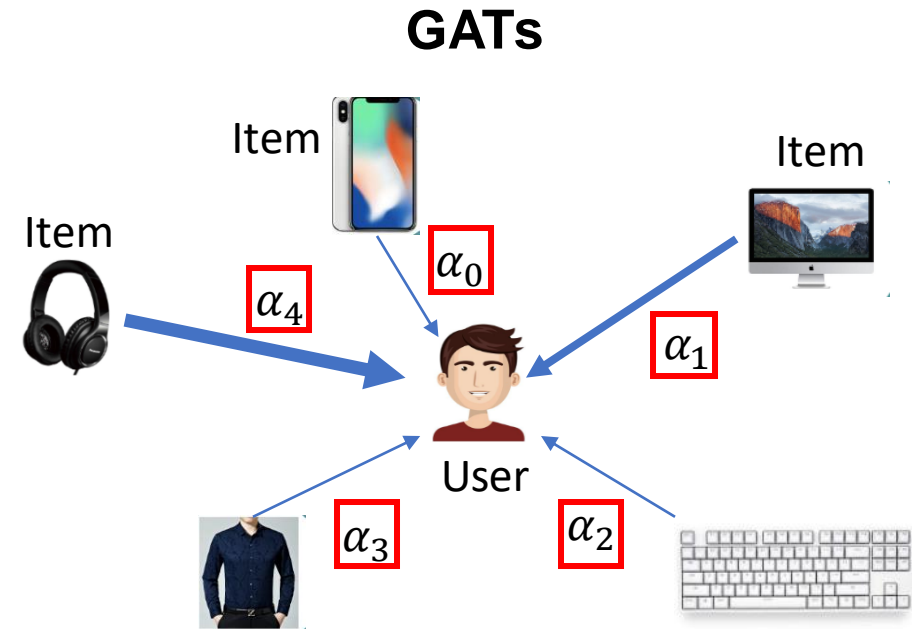


Looks similar to self-attention?

Transformers — in the Language of Graphs (1)



Step ① Mapping: Each node feature x_i is projected to q_i, k_i, v_i .




Attention computation: calculate the importance of neighbors

$$\alpha_{vu} = att(\mathbf{h}_v^{(l-1)}, \mathbf{h}_u^{(l-1)})$$

Transformers — in the Language of Graphs (2)

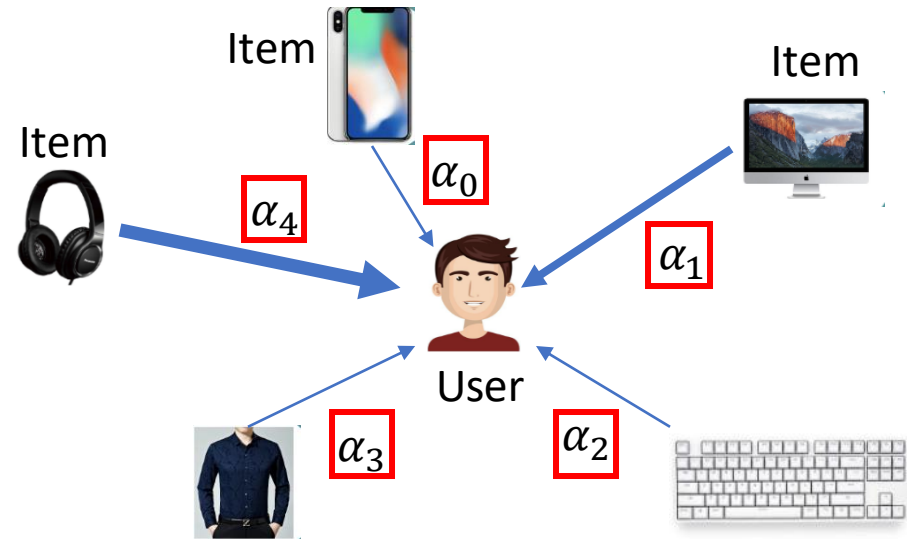
Transformers



Step ② Attention: Calculate the edge weights using $\mathbf{q}_i, \mathbf{k}_j$ of the two endpoints node i and j as $e_{ij} = \mathbf{q}_i^T \mathbf{k}_j / \sqrt{d}$, then normalizing it by the neighbors of node i

$$\alpha_{ij} = \text{softmax}_i(e_{ij}) = \frac{\exp(e_{ij})}{\sum_{k \in N_i} \exp(e_{ik})}$$

GATs




Message computing: transform information of neighbor node to a message


$$\mathbf{m}_u^{(l)} = \alpha_{vu} \mathbf{W}^{(l)} \mathbf{h}_u^{(l-1)}, u \in N_v$$

Transformers — in the Language of Graphs (3)

Transformers



GATs



Step ③ Update: Update each node feature according to its neighbors as

$$\mathbf{x}_i' = \sum_{k \in N_i} \alpha_{ik} \mathbf{x}_k$$

Aggregate message: aggregate messages from neighbor nodes

$$\mathbf{h}_v^{(l)} = \sigma \left(\sum_{u \in N_v} \mathbf{m}_u^{(l)} \right)$$

Transformers — in the Language of Graphs (4)

Summary: Comparison of **Self-attention (SA)** and **Graph Attention Networks (GAT)**

- Step ① Mapping
 - **SA**: different weights for q, k, v . $q = w_q x, k = w_k x, v = w_v x$.
 - **GAT**: shared weights for q, k, v . $q = wx, k = wx, v = wx$.
- Step ② Attention: **SA** uses dot-product attention, while (the original) **GAT** uses concatenation with MLP
 - Dot-product: $e_{ij} = q_i^T k_j / \sqrt{d}$
 - Concat: $e_{ij} = \text{act}(W [q_i || k_j])$, where c is a weight vector and act is the activation function like LeakyReLU

Transformers — in the Language of Graphs (5)

- The above computations do not require the assumption of **the complete graph**.
 - We assume full connectivity, mostly because we do not want to miss any potential token correlations.
- Self-attention can be easily adapted to graph-structured input data where the token correlations are given by the **adjacency matrix**, by replacing the **complete graph** with the **input graph**.
 - $\text{Self-Att}(X) = \text{Softmax}\left(\frac{(\mathbf{W}_k X)(\mathbf{W}_q X)^T}{\sqrt{d}} \odot \mathbf{A}_G \odot \mathbf{W}_E E\right) V$.
 - \mathbf{A}_G is the adjacency matrix of the graph and E is the edge weights of the graph if any.
- The complexity is no longer $O(n^2 d)$ but related to the edge number $O(E)$

Outline of Today's Lecture

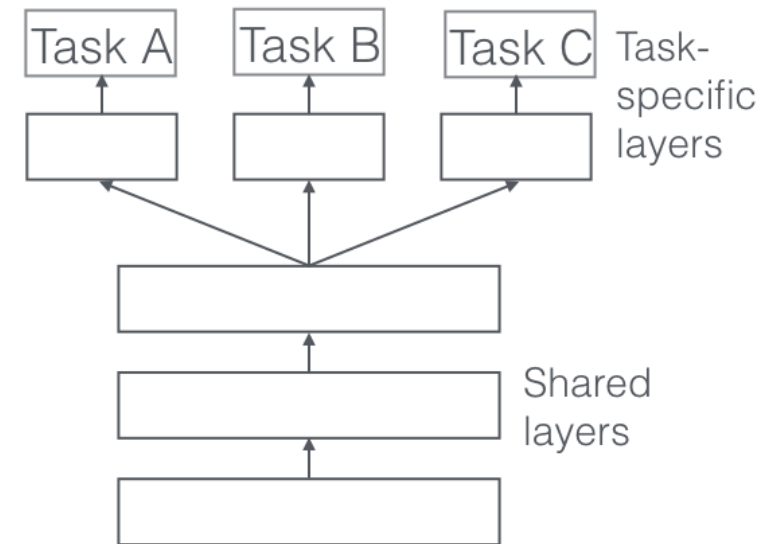
1. Self-Attention and Transformers

2. Transformers Applications

3. Graph Transformers and Sparse Transformers

Why is Transformer a Popular Choice

- Resolves various challenges of RNN-based architectures
- Attention makes the architecture **expressive and flexible** for different application scenarios
- It is very amenable to **self-supervised objectives**
 - We can leverage the vast number of **unsupervised examples** to learn a general model
 - Can be fine-tuned for **many downstream tasks**
 - Can out-perform models that are only trained for a specific downstream tasks

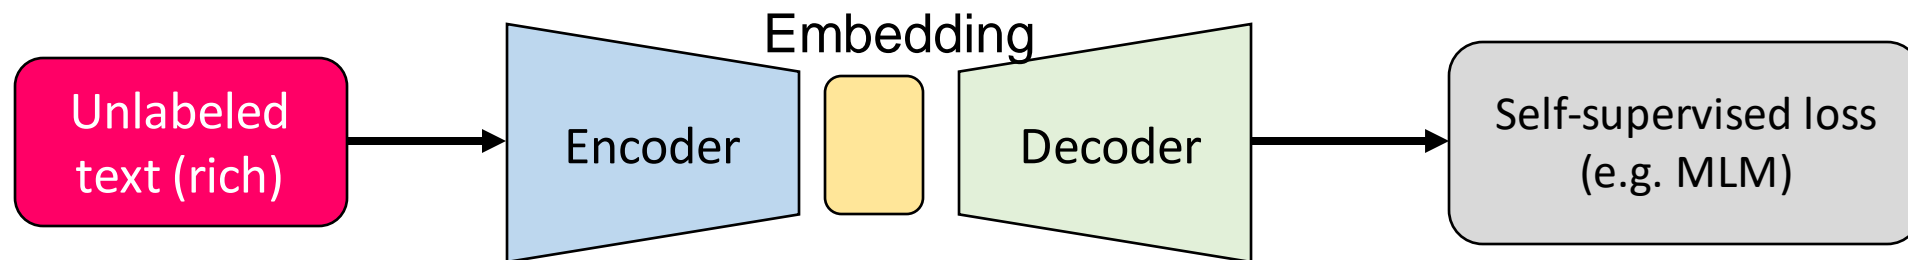


Label Scarcity

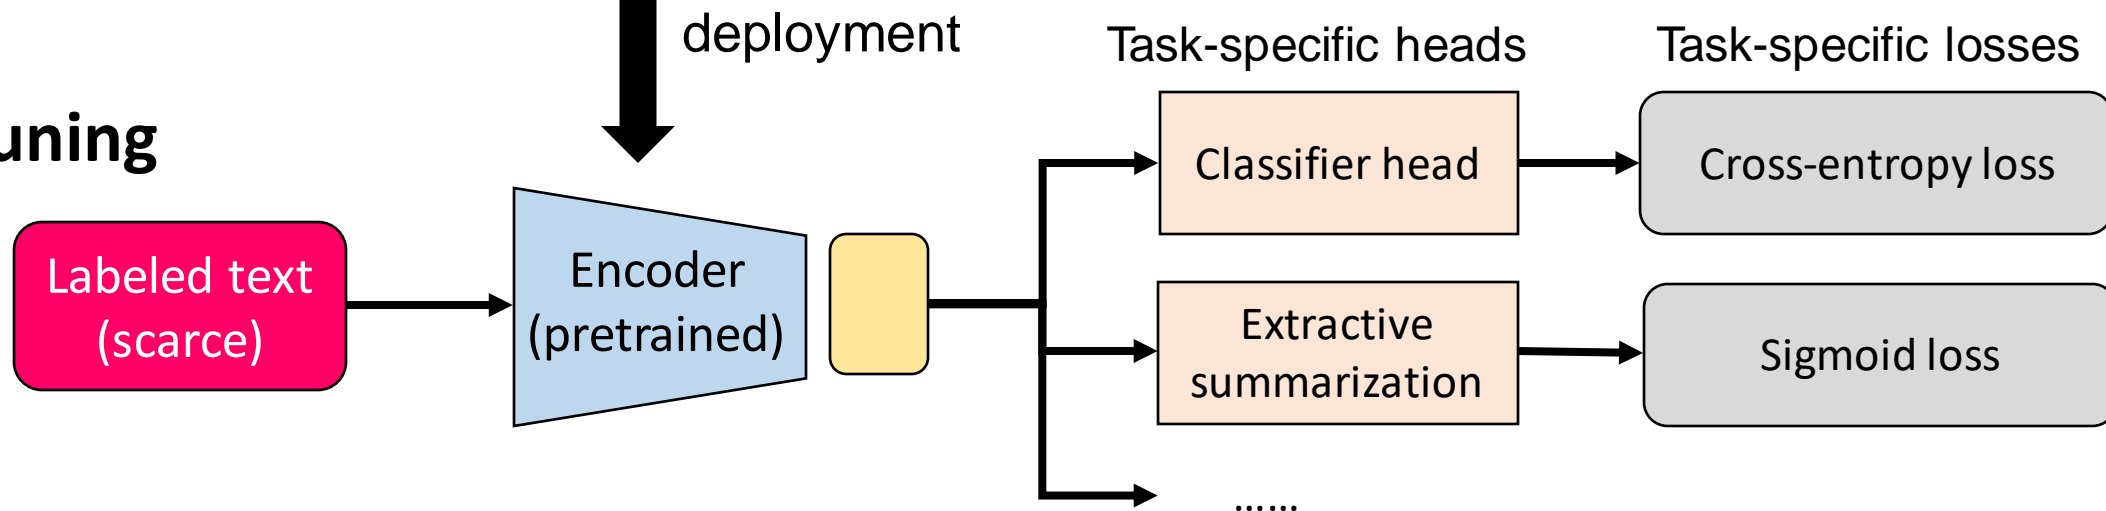
- ML models are hungry to data, especially labeled data for supervised task.
- The fast development of computer vision is largely contributed to **ImageNet**. It contains 14 million images **hand-annotated** by a team of researchers.
- This is often not possible for many domains. Most of time, it's easy to collect rich unlabeled data, but hard to obtain labeled data.
- Lack of annotated training example in NLP? **Pre-training** general-purpose language model on unlabeled large corpora (billions of characters) in **unsupervised** or **self-supervised** setting, then **fine-tuning** on smaller-scale tasks.

Pre-training and Fine-tuning

- **Pre-training**



- **Fine-tuning**




Transformers in NLP — BERT/RoBERTa (1)

BERT — **Bidirectional** Encoder Representations from Transformers [Devlin et al., 2018]

- **Pre-training task** (unsupervised): **Masked Language Model (MLM)**
 - First randomly masking $m\%$ tokens in the input sequence.
 - In BERT, 15% tokens are masked at random (replaced with the special [MASK] token)
 - Predicting masked tokens using remaining tokens.
 - Two modes: **Unidirectional** and **Bidirectional**.

Unidirectional [Radford et al., 2018]



- Maximize Likelihood of “all” given “Att” and “is”

Bidirectional




- Maximize Likelihood of “all”, given “Att”, “is”, “you”, “need”.

Transformers in NLP — BERT/RoBERTa (2)

BERT — **Bidirectional** Encoder Representations from Transformers [Devlin et al., 2018]


- **Masked Language Model (MLM)** loss function

Unidirectional



- Maximize Likelihood of “all” given “Att” and “is”
- $\max_{\theta} L(X) = \max_{\theta} \sum_{i \in M} \log P(x_i | x_{i-k}, \dots, x_{i-1}; \theta)$, where M is the set of tokens masked.
- In implementation,
 - Masking the upper triangle of attention matrix Att
 - $\text{Loss} = \sum_{i \in M} \text{CrossEntropy}(\text{Logit}_i, \text{Label}_i)$

Bidirectional



- Maximize Likelihood of “all”, given “Att”, “is”, “you”, “need”.
- $\max_{\theta} L(X) = \max_{\theta} \sum_{i \in M} \log P(x_i | \{x_j, i \neq j\}; \theta)$
- In implementation,
 - Keep full attention matrix Att
 - $\text{Loss} = \sum_{i \in M} \text{CrossEntropy}(\text{Logit}_i, \text{Label}_i)$

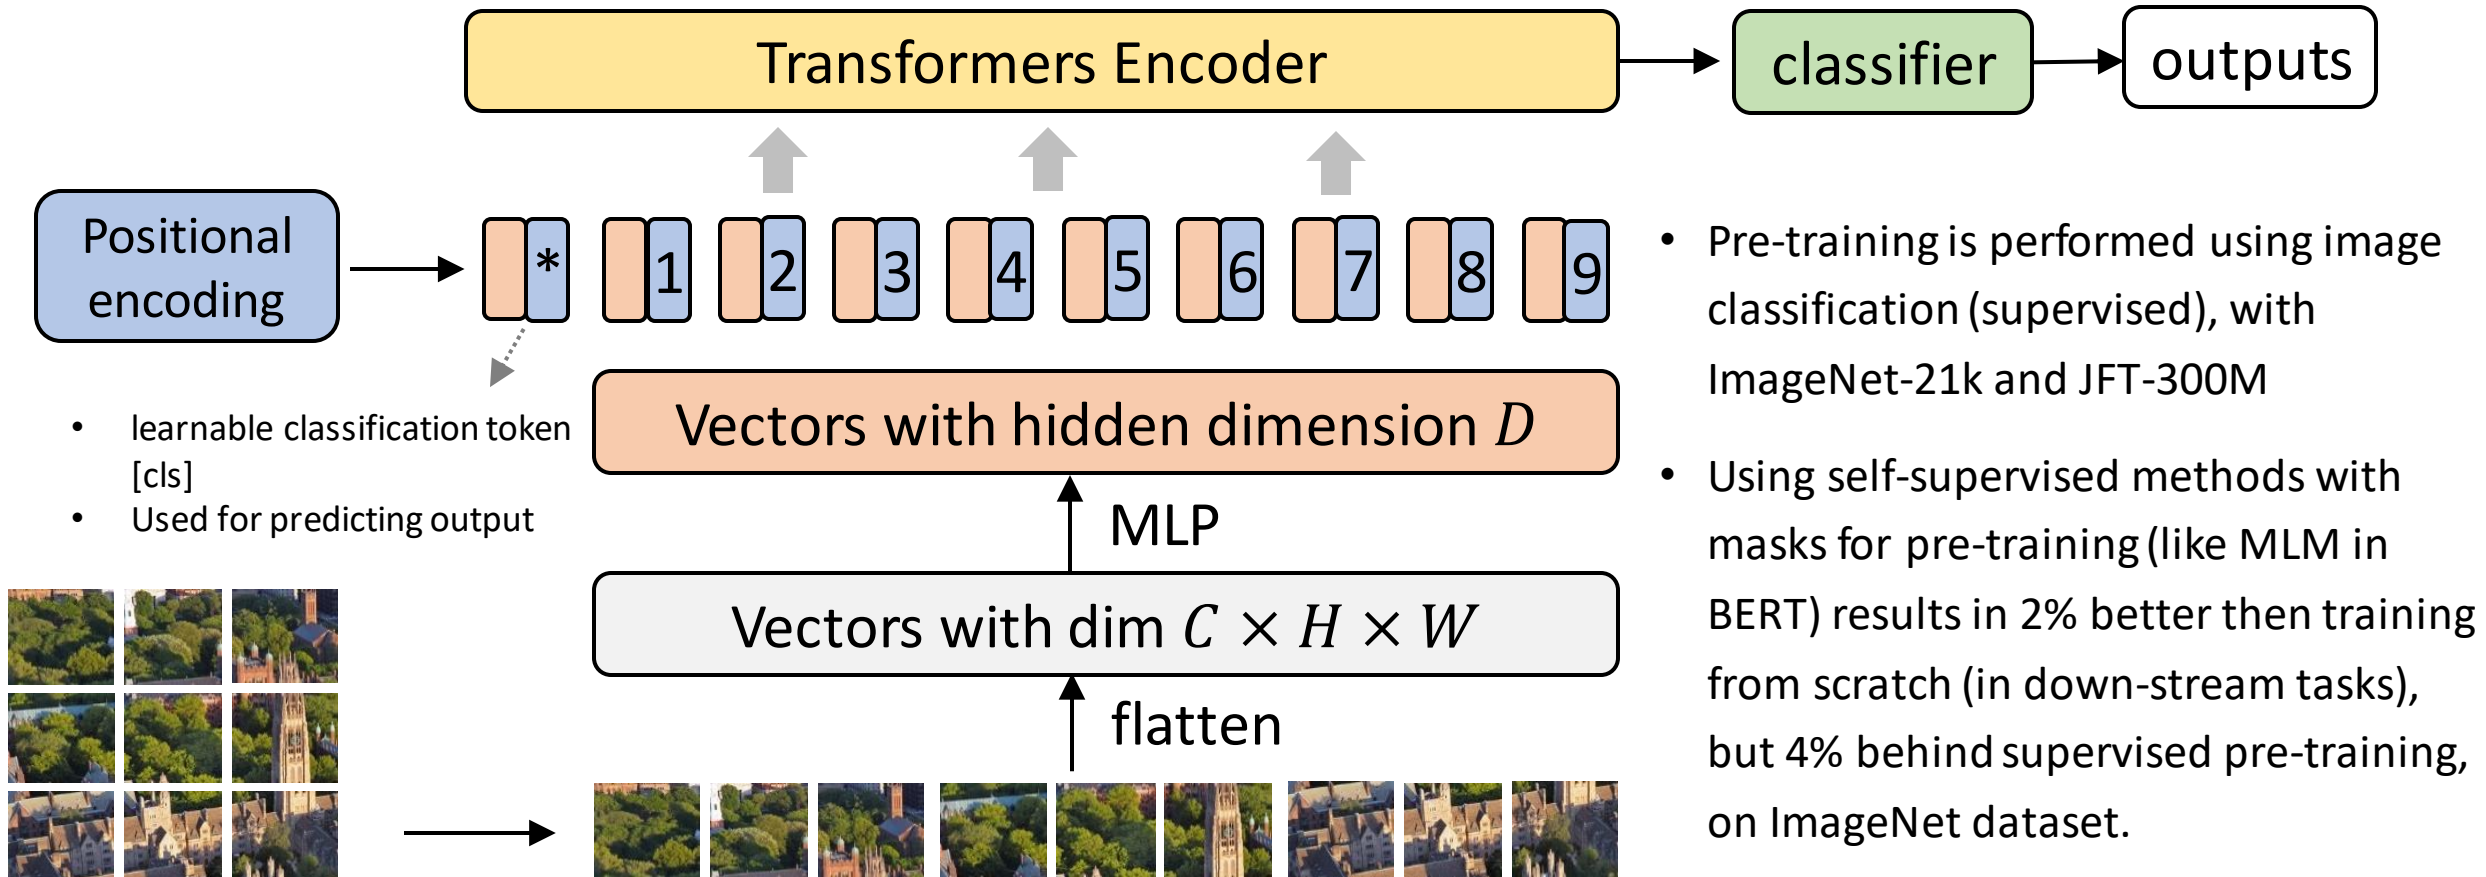
Transformers in NLP — BERT/RoBERTa (3)

RoBERTa — Robustly Optimized BERT [Liu et al., 2019]

- **Pretraining data**: BooksCorpus (800 M words) [Zhu et al., 2015], English Wiki (2500 M words), CC-News, OpenWebText [Gokaslan and Cohen, 2019], Stories [Trinh and Le, 2018]
 - Partition the corpus into “sentences” with fixed length of 512 tokens.
- **Hyperparameters** in use (also commonly adopted in most NLP Transformers):
 - **12-Layer** Encoder + **12-Layer** Decoder
(Pretrained Encoder is used more frequently in down-stream tasks)
 - Hidden dimension **768** = 12 (num of Heads) \times 64 (dim of Head)
 - Learning rate: Warmup then linear decay
 - Warmup: Gradually increasing the learning rate to a specific value in the first few epochs
 - Linear decay: Decreasing the learning rate by the same amount (decrement) every epoch.

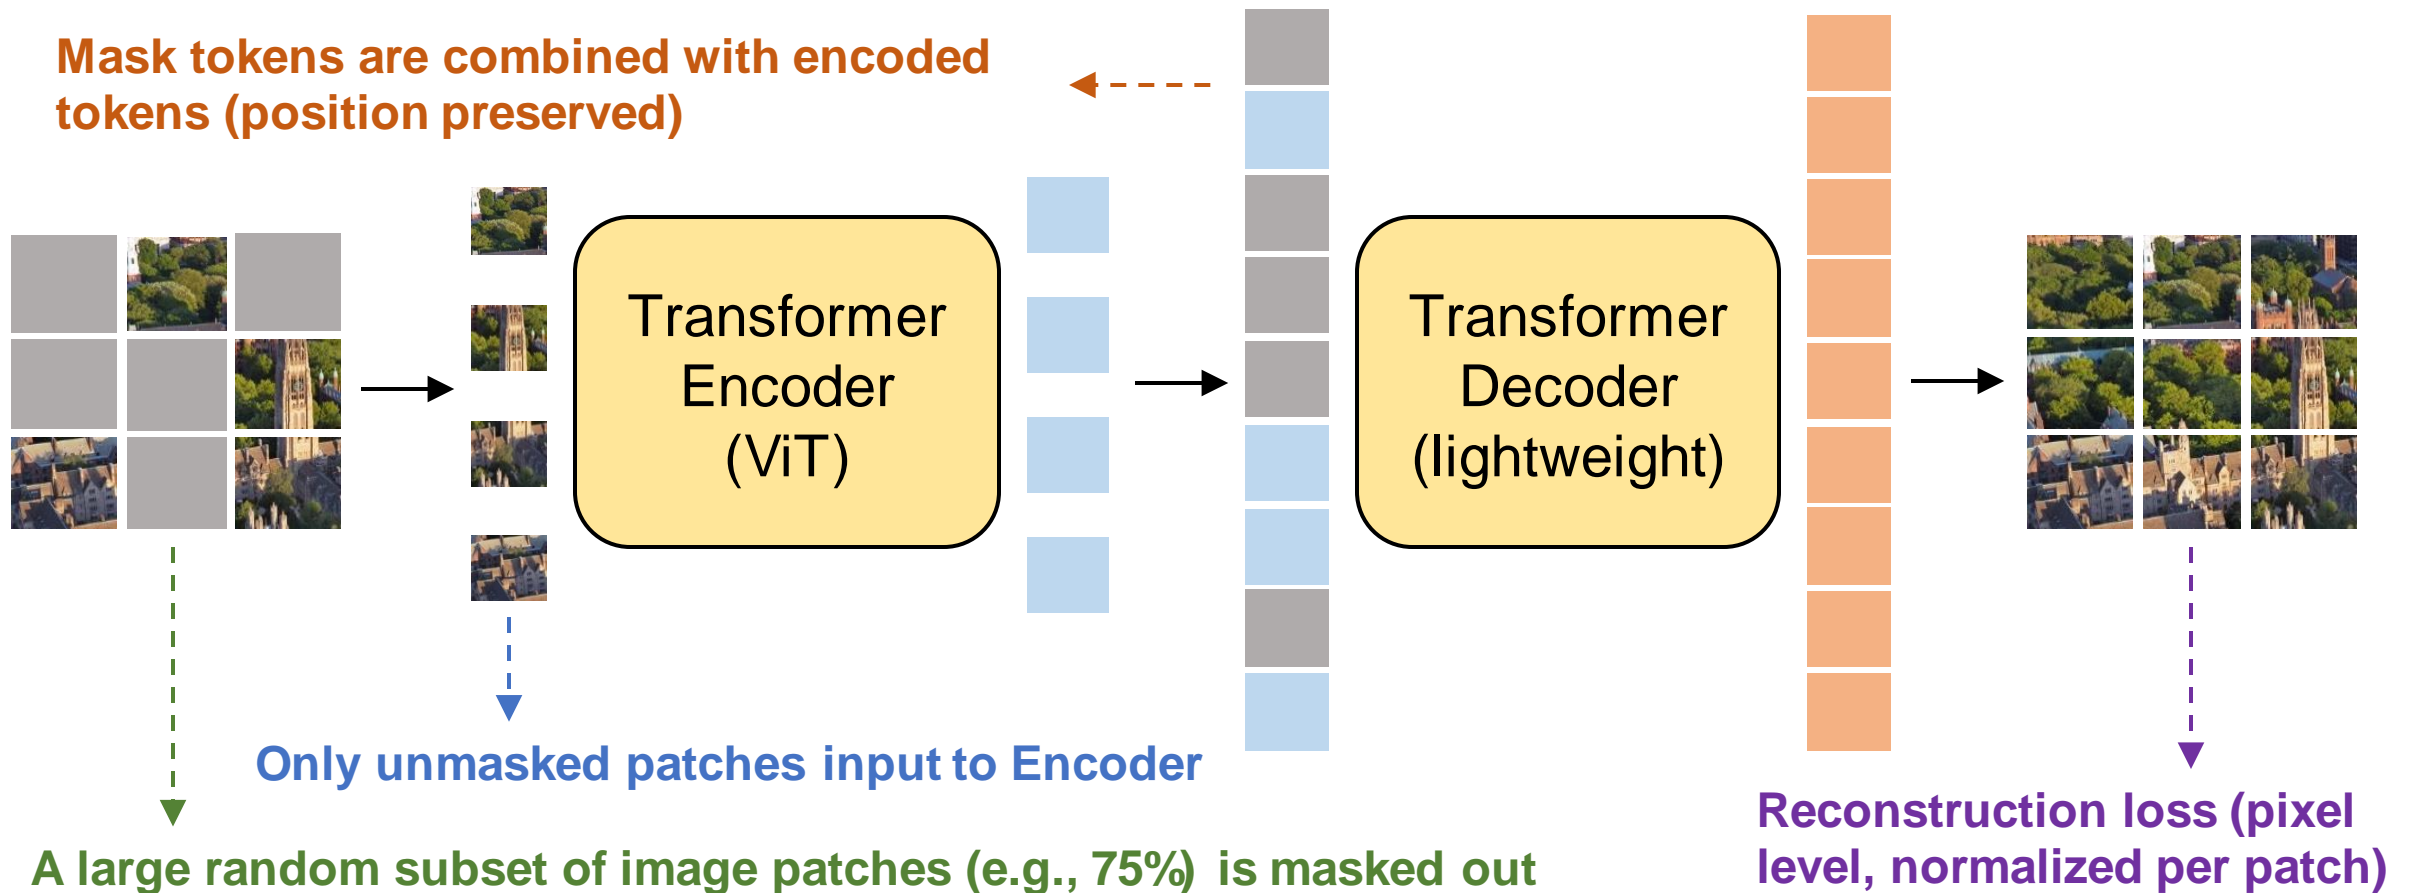
Transformers in CV — ViT [Dosovitskiy et al., ICLR 2021]

- An image patch is treated as a word in this context, and an image is partitioned to 16×16 tokens.




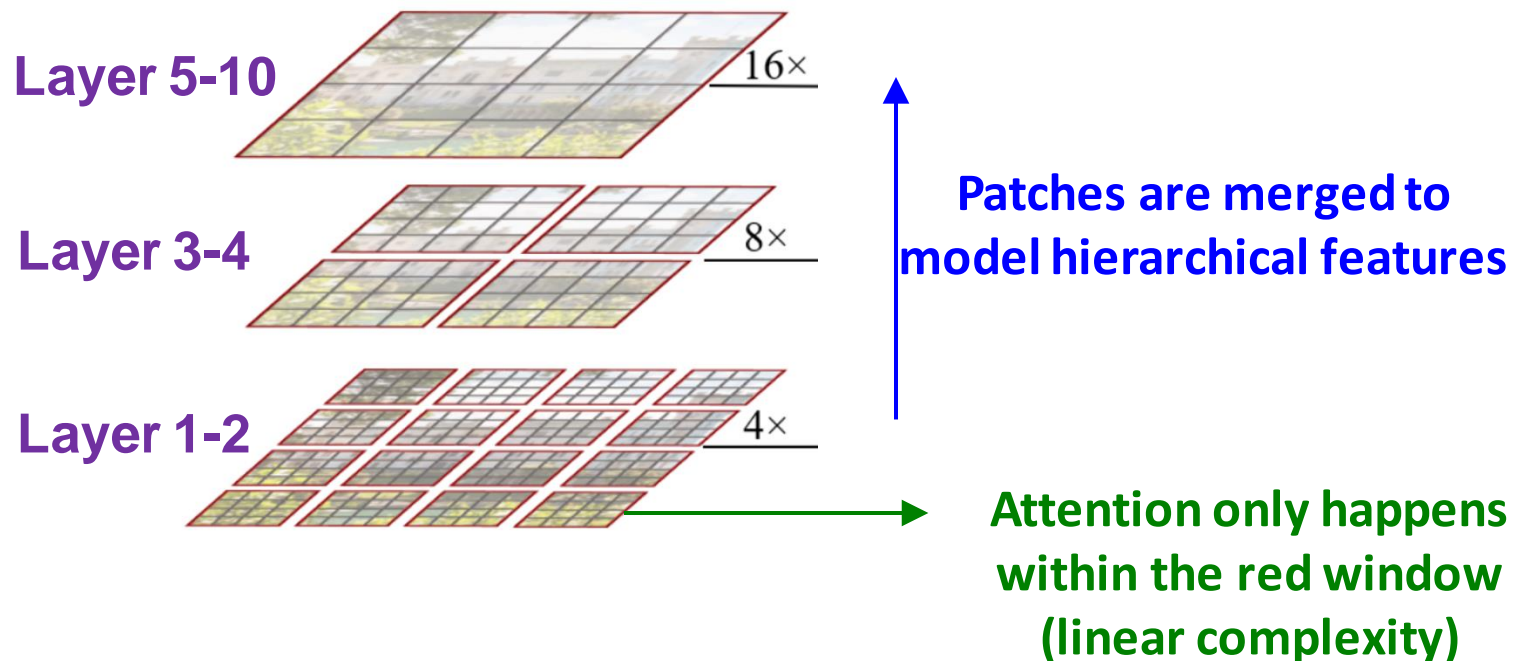
Transformers in CV — MAE [He et al., 2021]

- Can we use self-supervised pretraining for vision Transformers?
 - **Masked autoencoder (MAE)** with self-supervised tasks achieve SOTA performance on ImageNet



Transformers in CV — Swin [Liu et al., CVPR 2021]

- **Limitation of ViT**: using patch ($64 \times 64 \times 3$) as the unit token can be limited due to high resolution and variance of pixel values.
- **Solution**: Swin calculates fine-grained attention with hierarchical attention and shifted window.



Transformer Application Summary

- Transformer architectures are dominant when it comes to **self-supervised learning and pre-training**
 - The paradigm of pre-training and then fine-tuning / multi-task learning on downstream tasks has been very successful in many machine learning areas
 - **NLP** transformer (e.g. BERT)
 - **Vision** transformer (e.g. ViT)

Outline of Today's Lecture

1. Self-Attention and Transformers

2. Transformers Applications


3. Graph Transformers and Sparse Transformers

Transformers help Graph

Graphs help Transformers


Recall: Limitation of Single Hop Attention (1)

- A single GAT layer can only explore the relationship between node and its **one-hop** neighbors
 - Target node only attends to its immediate neighbors




Recall: Limitation of Single Hop Attention (2)

- A single hop attention falls short in exploring **broader graph structure** and **multi-hop** neighbors
 - Stacking multiple GAT layers causes **over-smoothing** and **over-fitting**




Multi-hop Attention v.s. Transformers (1)

Multi-hop Attention



- Enabling attention to **multi-hop** neighbors, resulting in direct/indirect attention between every pair of nodes


Transformers




- Enabling attention of target node to **all remaining nodes** in the graph, resulting in larger receptive field

Multi-hop Attention v.s. Transformers (2)

Multi-hop Attention



Transformers



- Using diffusion to reduce the complexity to $\mathcal{O}(E)$

- Full-attention with $\mathcal{O}(n^2)$ complexity

Graph Transformer as Matrix Operation

- No formal definitions of **Graph Transformers**
 - Self-attention can be defined among **1-hop neighbors**, **multi-hop neighbors**, or **all nodes** in the graph.
- E.g., **Graph Transformers** with **1-hop Attention**
 - In the matrix form
 - Self-Att(X) = $\text{softmax}\left(\frac{(\mathbf{W}_k X)(\mathbf{W}_q X)^T}{\sqrt{d}} \odot \mathbf{A}_G \odot \mathbf{W}_E \mathbf{E}\right) V$.
 - \mathbf{A}_G is the adjacency matrix of the graph and \mathbf{E} is the edge weights of the graph if any.
 - The complexity is no longer $O(n^2 d)$ but related to the edge number $O(E)$
- But this is not enough, any new challenges for graph Transformers?

\odot : Hadamard product (element-wise product)

Graph Transformers Challenges (1)

- **Challenge 1:** How to describe the position of a node in a **graph**?
 - **Permutation equivariance** should be preserved (this is why naïve one-hot PE does not work).
 - Weisfeiler-Lehman-PE [Zhang et al., Graph-BERT, arXiv 2020]; Laplacian PE [Dwivedi, arXiv 2020]
 - Random features (RF) or deterministic distance encoding (DE) [Li & Leskovec, 2021]
 - Besides graph Transformers, such PE are also applied in other GNNs to increase **expressiveness** (recall 1-WL constraints). Will discuss in future lectures.

Graph Transformers Challenges (2)

- **Challenge 2:** How to incorporate graph features (topology, edge weights...)
 - Naïve approach: restricting a node only attending to its neighbors
 - Graphormer [NeurIPS 2021]:
 - Attending to all nodes, besides its neighbors, $O(n^2)$ complexity

$$e_{ij} = \frac{(\mathbf{h}_i \mathbf{W}_q)(\mathbf{h}_j \mathbf{W}_k)^T}{\sqrt{d}} + b \phi(v_i, v_j) + c_{ij}$$

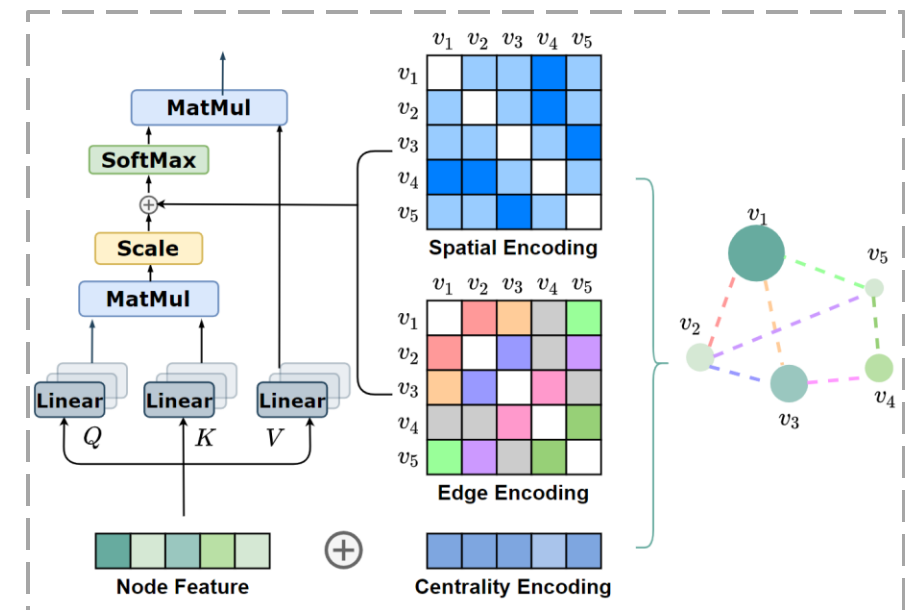
Spatial Encoding:

Shortest path between v_i, v_j

Edge Encoding:

Average all edge features along the shortest

path between v_i, v_j $c_{ij} = \frac{1}{N} \sum_{e \in SP(i,j)} x_e w_e$



Outline of Today's Lecture

1. Self-Attention and Transformers

2. Transformers Applications

3. Graph Transformers and Sparse Transformers


Transformers help Graph

Graphs help Transformers

Sparse Transformers

Conventional Transformers cannot scale to **long sequences** due to $O(n^2)$ complexity from the full-attention


- The QK^T matrix multiplication, Softmax(), AttV value updates all consume n^2 time and memory.
- Recall that Roberta and ViT are only designed for 512 (words) and 256 (patches), respectively.



Sparse Transformers

Conventional Transformers cannot scale to **long sequences** due to $O(n^2)$ complexity from the full-attention

- This is not enough for long sequences like a paragraph with **thousands** of words or an image with **$64 \times 64 \times 3 = 4096$** pixels.
- Because of quadratic dependency, increasing sequence length from 512 to 4096 results in around $64 \times$ more memory usage and running time, which usually cannot fit into a standard modern GPU (e.g NVIDIA A6000, 48 GB).

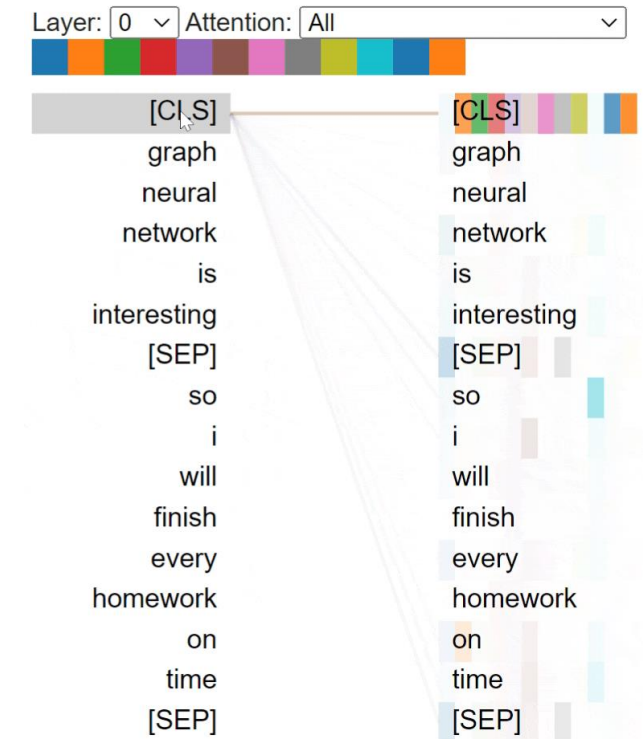


Sparse Transformers

- **Observation 1:** Although every attention is calculated, most of them are close to 0, the resulting attention maps are usually **sparse**.
- **Observation 2:** non-zero attention mostly appear between the node and its local neighbors. (**local** attention).
- **Observation 3:** some key words like “so” almost attend to every token in the sentence. (**global** attention)

Can we simplify self-attention (full-attention) using graph?

A sentence encoded by pretrained BERT

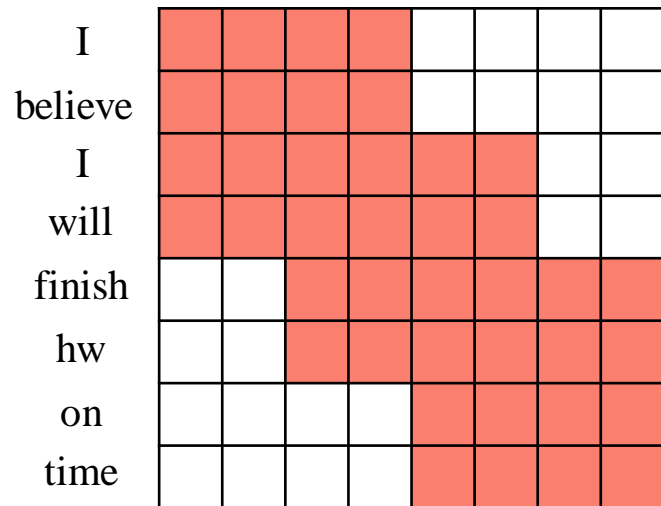


Try yourself!


<https://github.com/jessevig/bertviz#-quick-tour>

Sparse Transformers — Longformer [Beltagy et al., 2020]

- Applying overlapped local window attention to approximate the full-attention, only calculating attentions shaded in red



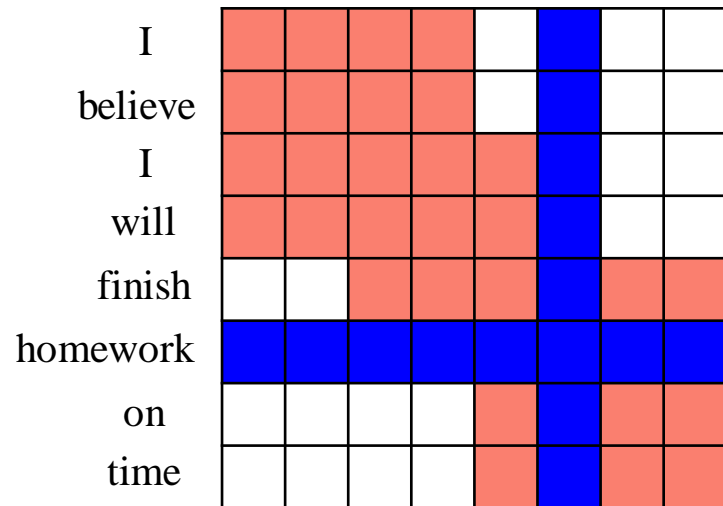
Masked Attention Pattern



Associated graph structure

Sparse Transformers — Longformer

- Longformer is based on the assumption that adjacent words have stronger correlations




Masked Attention Pattern

- Local window** is overlapped in half to enable cross-window attention (to ensure the graph is connected so that every pair of tokens can attend by stacking layers).
- Global attention** is further introduced for specific down-stream task
- Complexity is now $O(nw)$ compared to $O(n^2)$.
- Longformer can handle long sequences like 4096 tokens, by specifying local window size to be 512

Sparse Transformers — BigBird [Zaheer et al., 2020]

- BigBird model further introduces **Random Attention** to better approximate the full-attention.
- The smallest unit in BigBird is called a **block** (64 adjacent tokens)
- “Blockifying” is used to accelerate the sparse attention computation



Sparse Transformers — Reformer [Kitaev et al., 2020]

- Reformer is based on **locality sensitive hashing (LSH)** – a nearest-neighbor algorithm
 - **LSH** Definition: A hashing is called **locality-sensitive** if nearby vectors get the same hash with high probability and distant ones do no.
 - Main idea: First **grouping** tokens into different buckets (using LSH), then calculating **attentions within the same bucket**
 - The complexity is reduced to $O(b)$, b is the bucket size.


Sparse Transformers — Reformer [Kitaev et al., 2020]

- Details of **locality sensitive hashing (LSH)**

- To get b hashes (buckets), introducing a fixed random matrix $R \in \mathbb{R}^{d \times \frac{b}{2}}$, the hash of input token $\mathbf{x} \in \mathbb{R}^d$ is computed as

$$h(\mathbf{x}) = \operatorname{argmax}([\mathbf{x}R; -\mathbf{x}R])$$

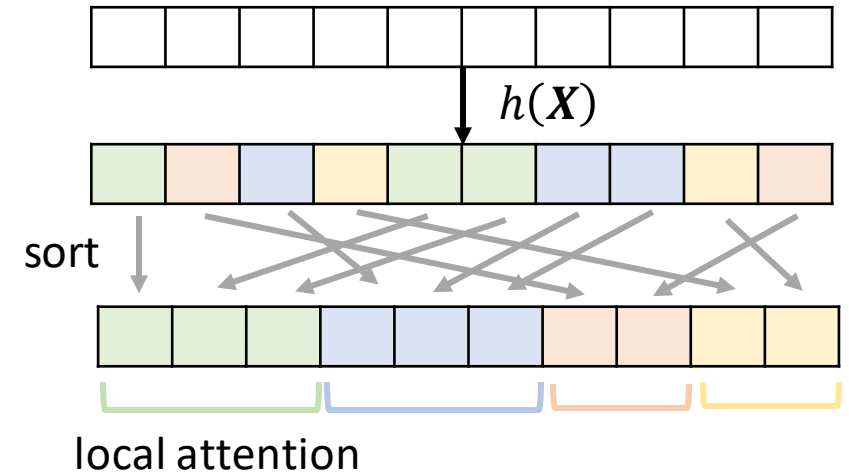

- Setting $d = 2$, \mathbf{x} becomes a 2-d vector
- Setting $b = 4$, we get 4 buckets (0, 1, 2, 3)
- For $\mathbf{x}' = R\mathbf{x}$, if we project \mathbf{x} and \mathbf{x}' to the unit circle, the function of $R \in \mathbb{R}^{2 \times 2}$ is rotation, i.e., $\mathbf{x} \xrightarrow{\text{rotated by } \theta} \mathbf{x}'$



$$h(\mathbf{x}) = \operatorname{argmax}([0.6, -0.8, -0.6, \mathbf{0.8}]) = \mathbf{3}$$

Sparse Transformers — Reformer [Kitaev et al., 2020]

- If two tokens \mathbf{x} , \mathbf{y} are close to each other (regarding e.g., angular distance), they have larger probabilities to share the same bucket.



- Attention in Reformer: only calculating attention within the same buckets

Summary

- Transformers
 - Transformers use **query (Q), key (K), value (V)** to compute attention values
 - All-pair correlation
 - Quadratic complexity
- Transformer applications
 - BERT (pre-trained **language** model)
 - ViT (**image** pre-training)
- (Multi-hop) graph attention is closely related to transformers!
 - Transformer can **inspire GNN architectures**
 - Graph learning techniques can **make transformers more efficient**