

# Distributed Node Embeddings

CPSC483: Deep Learning on Graph-Structured Data

Rex Ying

# Outline of Today's Lecture

- 1. Distributed Node Embeddings**
- 2. Random Walk Approaches for Node Embeddings**
- 3. Embedding Entire Graphs**

# Outline of Today's Lecture

## **1. Distributed Node Embeddings**

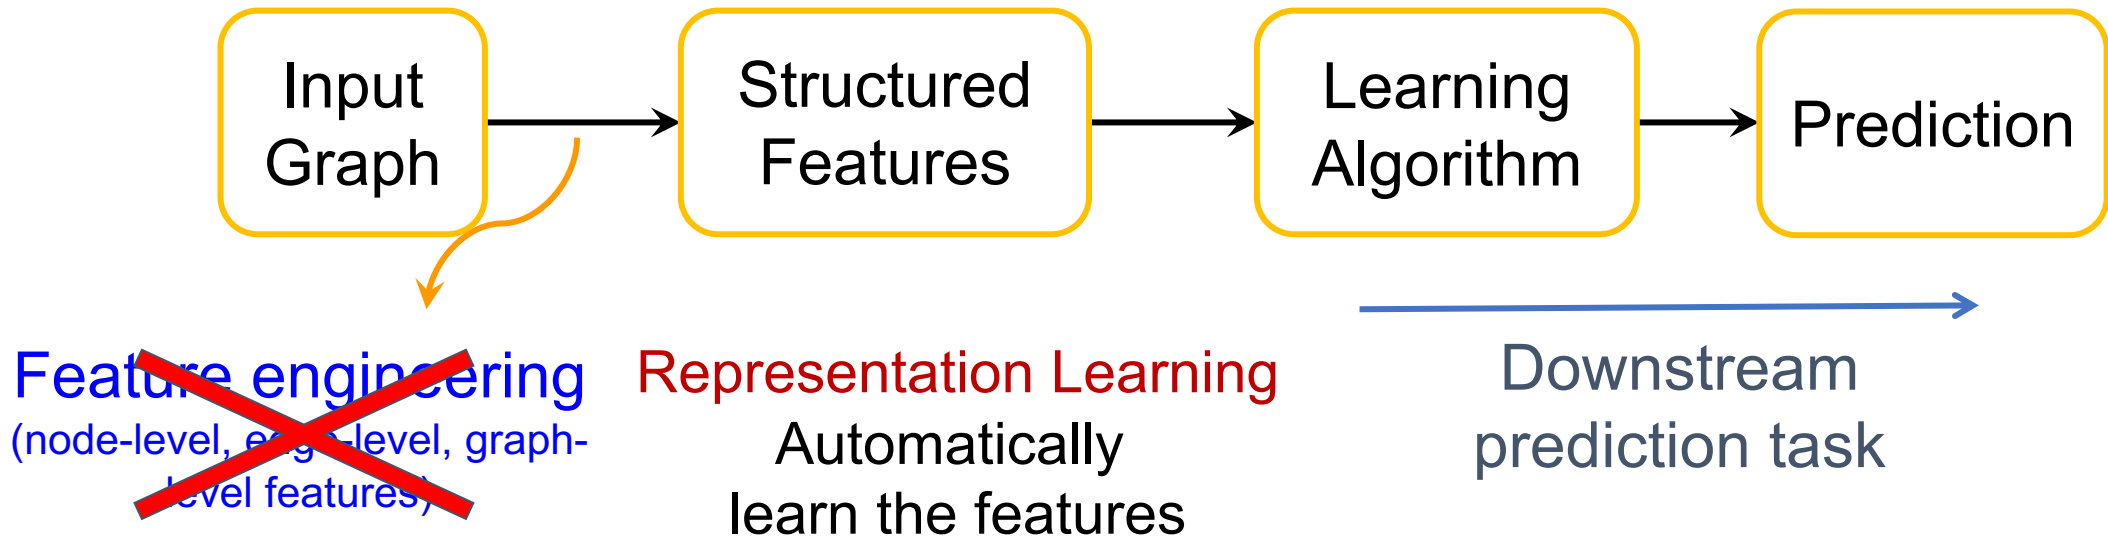
## 2. Random Walk Approaches for Node Embeddings

## 3. Embedding Entire Graphs

# Recap: Graph Representation Learning

In traditional machine learning, given an input graph, extract node, link and graph-level features, learn a model (SVM, neural network, etc.) that maps features to labels.

**Graph Representation Learning alleviates the need to do feature engineering **every single time**.**



# Recap: Learning Node Embeddings


An **unsupervised** setting for learning node embeddings

1. **Encoder ENC** maps from nodes to embeddings
2. **Define a node similarity function** (i.e., a measure of similarity in the original network)
3. **Decoder DEC** maps from embeddings to the similarity score
4. **Optimize the parameters of the encoder so that:**  
$$\text{similarity}(u, v) \approx \mathbf{z}_v^T \mathbf{z}_u$$
  
in the original network      Similarity of the embedding

# Recap: GNN for Node Embeddings

- graph neural networks (GNNs)  
encoder:

$\text{ENC}(\cdot)$  = multiple layers of  
non-linear transformations based  
on graph structures



**Today: “Shallow” Encoding!**

# “Shallow” Encoding (1)

- Consider a node  $v$  in a graph's nodes set  $\mathcal{V}$
- Simplest encoding approach: **Encoder is just an embedding-lookup**

embedding dimension

$$\text{ENC}(v) = \mathbf{z}_v = \mathbf{Z} \cdot v$$

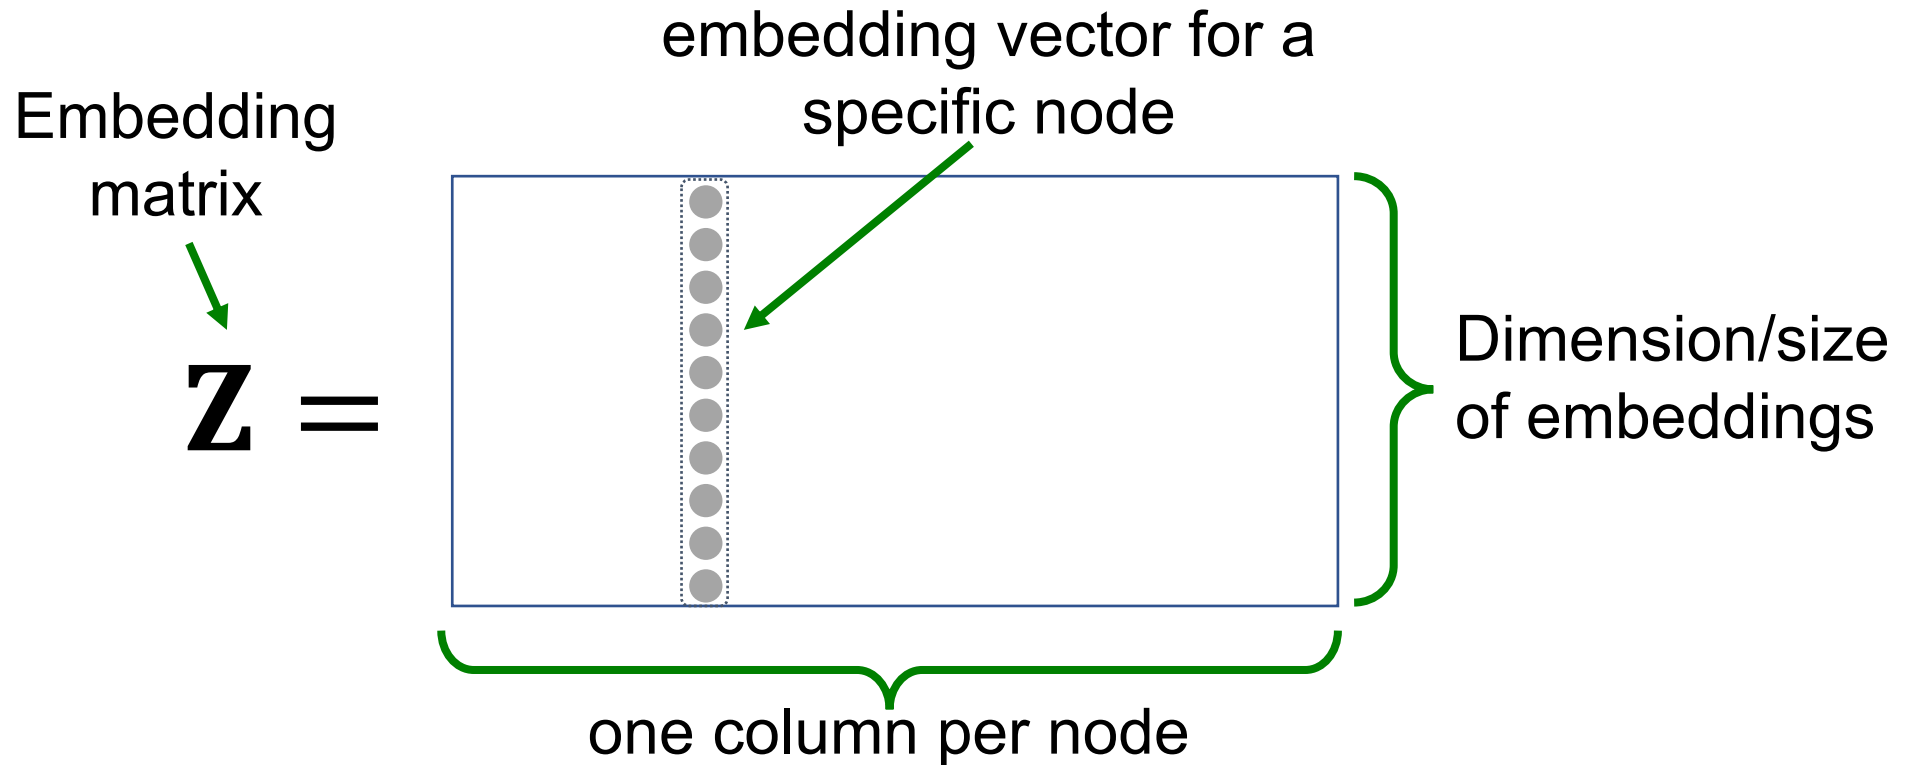
Encoder function

- $\mathbf{Z} \in \mathbb{R}^{d \times |\mathcal{V}|}$

- matrix, each column is a node embedding [what we learn / optimize]
- $v \in \mathbb{I}^{|\mathcal{V}|}$ 
  - Indicator vector, all zeroes except a one in column indicating node  $v$

# “Shallow” Encoding (2)

- Simplest encoding approach: **encoder is just an embedding-lookup**





# “Shallow” Encoding (3)

Simplest encoding approach: **Encoder is just an embedding-lookup**

**Each node is assigned a unique  
embedding vector**

(i.e., we directly optimize  
the embedding of each node)

Many methods: **DeepWalk, Node2Vec**

# Encoder + Decoder Framework Summary

- **Encoder + Decoder Framework**

- Shallow encoder: embedding **lookup**
- Parameters to optimize:  $\mathbf{Z}$  which contains node embeddings  $\mathbf{z}_u$  for all nodes  $u \in V$
- We will **not** cover deep encoders today.
- **Decoder:** based on node similarity.
- **Objective:** maximize  $\mathbf{z}_v^T \mathbf{z}_u$  for node pairs  $(u, v)$  that are **similar**

# How to Define Node Similarity

- Key choice of methods is **how they define node similarity**.
- Should two nodes have a similar embedding if they...
  - are linked?
  - share neighbors?
  - have similar “structural roles”?
- We will now learn node **similarity** definition that uses **random walks**, and how to optimize embeddings for such a similarity measure.

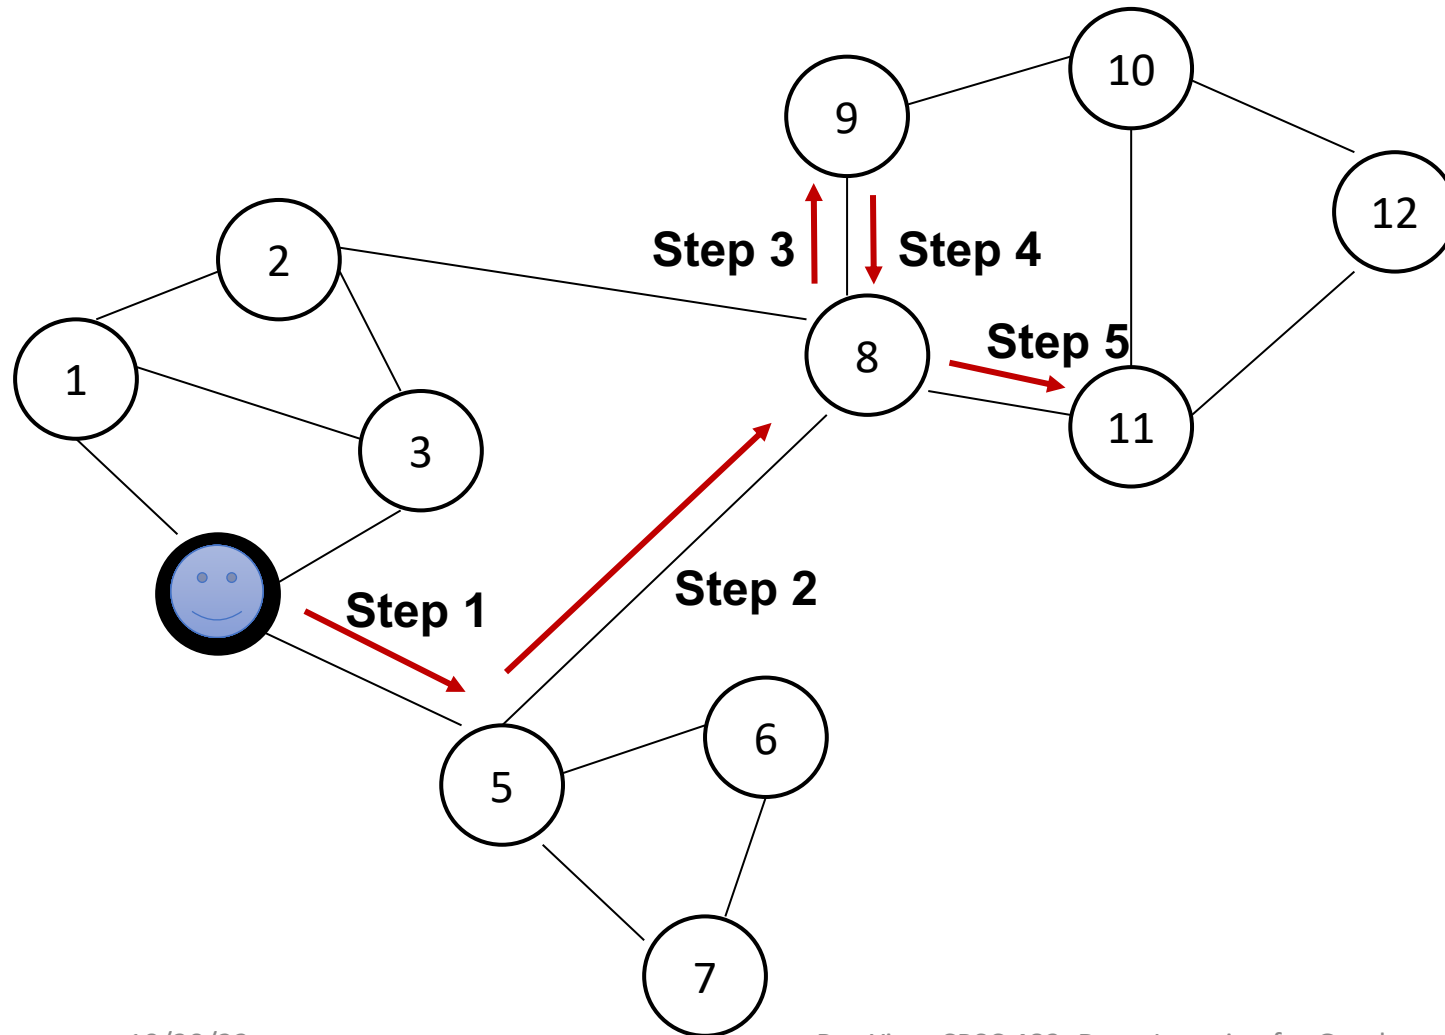
# Note on Distributed Node Embeddings

- This is **unsupervised/self-supervised** way of learning node embeddings
  - We are **not** utilizing node labels
  - We are **not** utilizing node features
  - The goal is to directly learn the embeddings of nodes so that some aspects of the network structure (captured by decoder) are preserved
- These embeddings are **task independent**
  - They are not trained for a specific task but can be used for any task.

# Outline of Today's Lecture

1. Non-GNN Node Embeddings
- 2. Random Walk Approaches for Node Embeddings**
3. Embedding Entire Graphs

# Random Walk




Given a *graph* and a *starting point*, we **select a neighbor** of it at **random**, and move to this neighbor; then we select a neighbor of this point at random, and move to it, etc. **The (random) sequence of points visited this way is a random walk on the graph.**

# Random Walk Embeddings (1)

$\mathbf{z}_u^T \mathbf{z}_v \approx$  probability that  $u$   
and  $v$  co-occur on  
a random walk  
over the graph


# Random Walk Embeddings (2)

1. Estimate probability of visiting node  $v$  on a random walk starting from node  $u$  using some random walk strategy  $R$



2. Optimize embeddings to encode these random walk statistics:

Similarity in embedding space (Here: dot product =  $\cos(\theta)$ ) encodes random walk “similarity”





# Why Random Walks?

1. **Expressivity:** Flexible stochastic definition of node similarity that incorporates both local and higher-order neighborhood information  
**Idea:** if random walk starting from node  $u$  visits  $v$  with high probability,  $u$  and  $v$  are similar (high-order multi-hop information)
2. **Efficiency:** Do not need to consider all node pairs when training; only need to consider pairs that co-occur on random walks

# Unsupervised Feature Learning

- **Intuition:** Find embedding of nodes in  $d$ -dimensional space that preserves similarity
- **Idea:** Learn node embedding such that **nearby** nodes are close together in the network
- **Given a node  $u$ , how do we define nearby nodes?**
  - $N_R(u)$  ... neighbourhood of  $u$  obtained by some random walk strategy  $R$

# Feature Learning as Optimization

- Given  $G = (V, E)$ ,
  - Our goal is to learn a mapping  $f: u \rightarrow \mathbb{R}^d$ :  
 $f(u) = \mathbf{z}_u$

- Log-likelihood objective:

$$\max_f \sum_{u \in V} \log P(N_R(u) | \mathbf{z}_u)$$

- $N_R(u)$  is the neighborhood of node  $u$  by strategy  $R$
- Given node  $u$ , we want to learn feature representations that are predictive of the nodes in its random walk neighborhood  $N_R(u)$

# Random Walk Optimization (1)

1. Run **short fixed-length random walks** starting from each node  $u$  in the graph using some random walk strategy  $R$
2. For each node  $u$  collect  $N_R(u)$ , the multiset\* of nodes visited on random walks starting from  $u$
3. Optimize embeddings according to: **Given node  $u$ , predict its neighbors  $N_R(u)$**

$$\max_f \sum_{u \in V} \log P(N_R(u) | \mathbf{z}_u) \quad \Rightarrow \quad \text{Maximum likelihood objective}$$

\* $N_R(u)$  can have repeat elements since nodes can be visited multiple times on random walks

# Random Walk Optimization (2)

- Equivalently,

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log(P(v|\mathbf{z}_u))$$

- **Intuition:** Optimize embeddings  $\mathbf{z}_u$  to maximize the likelihood of random walk co-occurrences
- **Parameterize  $P(v|\mathbf{z}_u)$  using softmax:**

$$P(v|\mathbf{z}_u) = \frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}$$

## Why softmax?

We want node  $v$  to be most similar to node  $u$  (out of all nodes  $n$ ).

**Intuition:**  $\sum_i \exp(x_i) \approx \max_i \exp(x_i)$

# Random Walk Optimization (3)

- Putting it all together

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} - \log \left( \frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right)$$

sum over all nodes  $u$       sum over nodes  $v$  seen on random walks starting from  $u$       predicted probability of  $u$  and  $v$  co-occurring on random walk

- Optimizing random walk embeddings = Finding embeddings  $\mathbf{z}_u$  that minimize  $\mathcal{L}$

# Random Walk Optimization (4)

- But doing this naively is too expensive!

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log\left(\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}\right)$$

Nested sum over nodes  
gives  $O(|V|^2)$  complexity!

# Random Walk Optimization (5)

- But doing this naively is too expensive!

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log\left(\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}\right)$$

Nested sum over nodes  
gives  $O(|V|^2)$  complexity!

The normalization term  
from the softmax is the  
culprit... can we  
approximate it?



# Negative Sampling (1)

- **Solution:** Negative Sampling

## Why is the approximation valid?

Technically, this is a different objective. But Negative Sampling is a form of Noise Contrastive Estimation (NCE) which approx. maximizes the log probability of softmax.

New formulation corresponds to using a logistic regression (sigmoid func.) to distinguish the target node  $v$  from nodes  $n_i$  sampled from background distribution  $P_v$ .

More at <https://arxiv.org/pdf/1402.3722.pdf>

$$\log \left( \frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right) \approx \log \left( \sigma(\mathbf{z}_u^T \mathbf{z}_v) \right) - \sum_{i=1}^k \log \left( \sigma(\mathbf{z}_u^T \mathbf{z}_{n_i}) \right), n_i \sim P_V$$

sigmoid function

Random distribution over nodes

Instead of normalizing w.r.t. all nodes, just normalize against  $k$  random “negative samples”  $n_i$

# Negative Sampling (2)

$$\log \left( \frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)} \right) \approx \log \left( \sigma(\mathbf{z}_u^T \mathbf{z}_v) \right) - \sum_{i=1}^k \log \left( \sigma(\mathbf{z}_u^T \mathbf{z}_{n_i}) \right), n_i \sim P_V$$

- Sample  $k$  negative nodes each with prob. proportional to its degree
  - Two consideration for  $k$  (# negative samples):
    - Higher  $k$  gives more robust estimates
    - Higher  $k$  corresponds to higher bias on negative events
- In practice,  $k = 5 \sim 20$

# Random Walk Optimization (6)

- After we obtained the objective function, how do we optimize (minimize) it?

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log(P(v|\mathbf{z}_u))$$

- **Solution:** Gradient Descent

- $\mathbf{z}_i \leftarrow \mathbf{z}_i - \eta \frac{\partial \mathcal{L}}{\partial \mathbf{z}_i}, i \in \mathcal{V}$

# Random Walk: Summary

1. Run short fixed-length random walks starting from each node on the graph
2. For each node  $u$  collect  $N_R(u)$ , the multiset of nodes visited on random walks starting from  $u$
3. Optimize embeddings using Stochastic Gradient Descent:

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log(P(v|\mathbf{z}_u))$$

**We can efficiently approximate this  
using negative sampling!**

# How should We Randomly Walk?


- So far we have described how to optimize embeddings given a random walk strategy  $R$
- **What strategies should we use to run these random walks?**
  - Simplest idea: **Just run fixed-length, unbiased random walks starting from each node** (i.e., [DeepWalk from Perozzi et al., 2013](#))
    - The issue is that such notion of similarity is too constrained
- **How can we generalize this?**

# Overview of Node2Vec

- **Goal:** Embed nodes with similar network neighborhoods close in the feature space.
- We frame this goal as a maximum likelihood optimization problem, independent to the downstream prediction task.
- **Key observation:** Flexible notion of network neighborhood  $N_R(u)$  of node  $u$  leads to rich node embeddings
- Develop biased 2<sup>nd</sup> order random walk  $R$  to generate network neighborhood  $N_R(u)$  of node  $u$


# Node2Vec: Biased Walks (1)

- **Idea:** use flexible, biased random walks that can trade off between **local** and **global** views of the network ([Grover and Leskovec, 2016](#)).



# Node2Vec: Biased Walks (2)


- Two classic strategies to define a neighborhood  $N_R(u)$  of a given node  $u$ :



- Walk of length 3 ( $N_R(u)$  of size 3):
  - $N_{BFS}(u) = \{s_1, s_2, s_3\}$  Local microscopic view
  - $N_{DFS}(u) = \{s_4, s_5, s_6\}$  Global macroscopic view




# BFS vs. DFS



**BFS:**

Micro-view of  
neighbourhood



**DFS:**

Macro-view of  
neighbourhood

# Interpolating BFS and DFS


**Biased fixed-length random walk  $R$  that given a node  $u$  generates neighborhood  $N_R(u)$**

- Two parameters:
  - **Return parameter  $p$ :**
    - Return back to the previous node
  - **In-out parameter  $q$ :**
    - Moving outwards (DFS) vs. inwards (BFS)
    - Intuitively,  $q$  is the “ratio” of BFS vs. DFS

# Biased Random Walks (1)

## Biased 2<sup>nd</sup>-order random walks explore network neighborhoods:


- Rnd. walk just traversed edge  $(s_1, w)$  and is now at  $w$
- **Insight:** Neighbors of  $w$  can only be:



**Idea:** Remember where the walk came from

# Biased Random Walks (2)

- Walker came over edge  $(s_1, w)$  and is at  $w$ . Where to go next?




$1/p, 1/q, 1$  are  
unnormalized  
probabilities

- $p, q$  model transition probabilities
  - $p$  ... return parameter
  - $q$  ... "walk away" parameter

# Biased Random Walks (3)

- Walker came over edge  $(s_1, w)$  and is at  $w$ . Where to go next?



Target $t$	Prob.	Dist. $(s_1, t)$
$s_1$	$1/p$	0
$s_2$	1	1
$s_3$	$1/q$	2
$s_4$	$1/q$	2

- BFS-like** walk: Low value of  $p$
- DFS-like** walk: Low value of  $q$

$N_R(u)$  are the nodes visited by the biased walk

Unnormalized  
transition prob.  
segmented based on  
distance from  $s_1$

# Node2Vec Algorithm

1. Compute random walk probabilities
  2. Simulate  $r$  random walks of length  $l$  starting from each node  $u$
  3. Optimize the node2vec objective using stochastic gradient descent
- Linear-time complexity
  - All 3 steps are individually parallelizable

# Other Random Walk Ideas

- **Different kinds of biased random walks:**
  - Based on node attributes ([Dong et al., 2017](#)).
  - Based on learned weights ([Abu-El-Haija et al., 2017](#))
- **Alternative optimization schemes:**
  - Directly optimize based on 1-hop and 2-hop random walk probabilities (as in [LINE from Tang et al. 2015](#)).
- **Network preprocessing techniques:**
  - Run random walks on modified versions of the original network (e.g., [Ribeiro et al. 2017's struct2vec](#), [Chen et al. 2016's HARP](#)).

# Summary of Part 2

- **Core idea:** Embed nodes so that distances in embedding space reflect node similarities in the original network.
- **Different notions of node similarity:**
  - Naïve: similar if 2 nodes are connected
  - Neighborhood overlap (covered in Lecture 2)
  - Random walk approaches (**covered today**)
- **So what method should I use..?**
- No one method wins in all cases....
  - E.g., node2vec performs better on node classification while alternative methods perform better on link prediction ([Goyal and Ferrara, 2017 survey](#))
- Random walk approaches are generally more efficient
- **In general:** Must choose definition of node similarity that matches your application!




# Outline of Today's Lecture

1. Non-GNN Node Embeddings
2. Random Walk Approaches for Node Embeddings
- 3. Embedding Entire Graphs**

# Embedding Entire Graphs


- **Goal:** Want to embed a subgraph or an entire graph  $G$ . Graph embedding:  $\mathbf{z}_G$ .



- **Tasks:**
  - Classifying toxic vs. non-toxic molecules
  - Identifying anomalous graphs

# Embedding Entire Graphs: Approaches

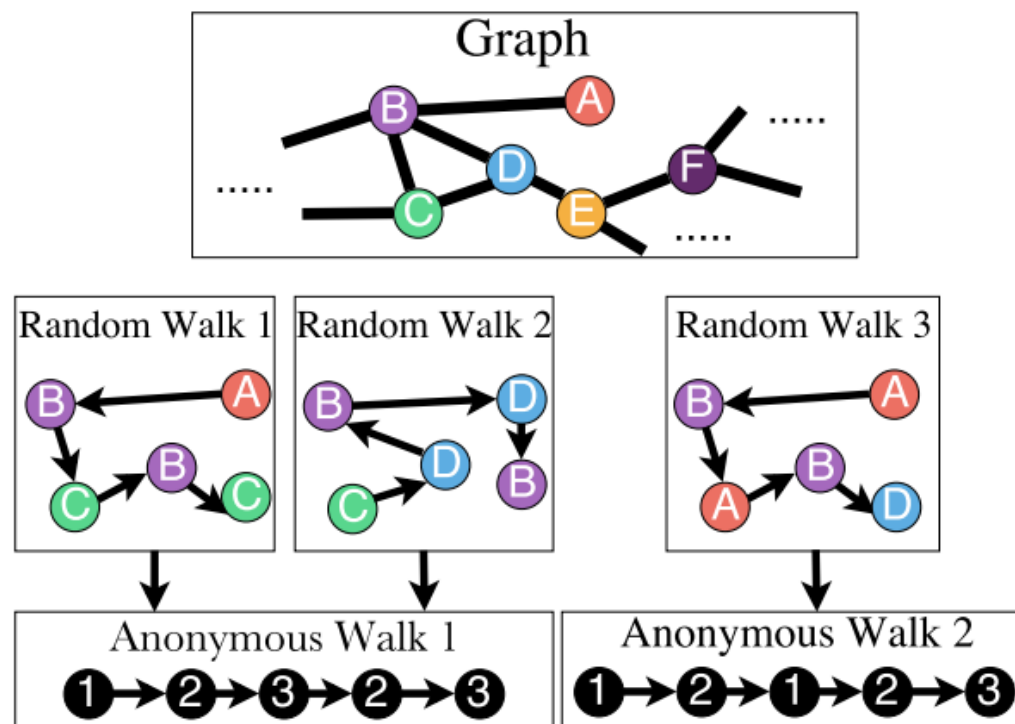
- **Approach 1:** (Recall in lecture 5) sum/mean/max/hierarchical pooling
- **Approach 2:** add **virtual node**
  - Proposed by [Li et al., 2016](#) as a general technique for subgraph embedding



- **Approach 3:** Anonymous walk embeddings

# Anonymous Walk Embeddings (1)



- States in **anonymous walks** correspond to the index of the **first time** we visited the node in a random walk



Anonymous Walk Embeddings, ICML 2018 <https://arxiv.org/pdf/1805.11921.pdf>

# Anonymous Walk Embeddings (2)

- Agnostic to the identity of the nodes visited (hence anonymous)
- Example RW1 (Random Walk 1):
  - Step 1: node A  $\rightarrow$  node 1
  - Step 2: node B  $\rightarrow$  node 2 (different from node 1)
  - Step 3: node C  $\rightarrow$  node 3 (different from node 1, 2)
  - Step 4: node B  $\rightarrow$  node 2 (same as the node in step 2)
  - Step 5: node C  $\rightarrow$  node 3 (same as the node in step 3)
- Note: RW2 gives the same anonymous walk




# Number of Walks Grows

Number of anonymous walks grows exponentially:

- There are 5 anon. walks  $w_i$  of length 3:

$w_1=111, w_2=112, w_3=121,$

$w_4=122, w_5=123$



# Simple Use of Anonymous Walks

- Simulate anonymous walks  $w_i$  of  $l$  steps and record their counts
- **Represent the graph as a probability distribution over these walks**
- **For example:**
  - Set  $l = 3$
  - Then we can represent the graph as a 5-dim vector
    - Since there are 5 anonymous walks  $w_i$  of length 3: 111, 112, 121, 122, 123
  - $\mathbf{Z}_G[i] = \text{probability of anonymous walk } w_i \text{ in } G$

# Sampling Anonymous Walks

- **Sampling anonymous walks:** Generate independently a set of  $m$  random walks
- Represent the graph as a probability distribution over these walks

- How many random walks  $m$  do we need?

- We want the distribution to have error of more than  $\varepsilon$  with prob. less than  $\delta$ :

$$m = \left\lceil \frac{2}{\varepsilon^2} (\log(2^\eta - 2) - \log(\delta)) \right\rceil$$

- where:  $\eta$  is the total number of anon. walks of length  $l$ .

**For example:**

There are  $\eta = 877$  anonymous walks of length  $l = 7$ . If we set  $\varepsilon = 0.1$  and  $\delta = 0.01$  then we need to generate  $m=122,500$  random walks



# New Idea: Learn Walk Embeddings

Rather than simply represent each walk by the fraction of times it occurs, we **learn embedding  $z_i$  of anonymous walk  $w_i$**

- Learn a graph embedding  $z_G$  together with all the anonymous walk embeddings  $z_i$

$Z = \{z_i : i = 1 \dots \eta\}$ , where  $\eta$  is the number of sampled anonymous walks.

- Note that  $z_i$  are embeddings of anonymous walk now instead of embeddings of nodes.



## How to embed walks?

- **Idea:** Embed walks s.t. the next walk can be predicted

Anonymous Walk Embeddings, ICML 2018 <https://arxiv.org/pdf/1805.11921.pdf>

# Learn Walk Embeddings (1)

- A vector parameter  $\mathbf{z}_G$  for input graph
  - The embedding of entire graph to be learned
- Starting from **node 1**: Sample anonymous random walks, e.g.



- **Learn to predict walks that co-occur in  $\Delta$ -size window** (e.g. predict  $w_2$  given  $w_1$ ,  $w_3$  if  $\Delta = 1$ )
- Objective:

$$\max \sum_{t=\Delta}^{T-\Delta} \log P(w_t | w_{t-\Delta}, \dots, w_{t+\Delta}, \mathbf{z}_G)$$

- Sum the objective over all nodes in the graph

# Learn Walk Embeddings (2)

- Run  $T$  different random walks from  $u$  each of length  $l$ :


$$N_R(u) = \{w_1^u, w_2^u \dots w_T^u\}$$

- Learn to predict walks that co-occur in  $\Delta$ -size window
- Estimate embedding  $\mathbf{z}_i$  of anonymous walk  $w_i$

**Objective:**  $\max \frac{1}{T} \sum_{t=\Delta}^{T-\Delta} \log P(w_t | \{w_{t-\Delta}, \dots, w_{t+\Delta}, \mathbf{z}_G\})$

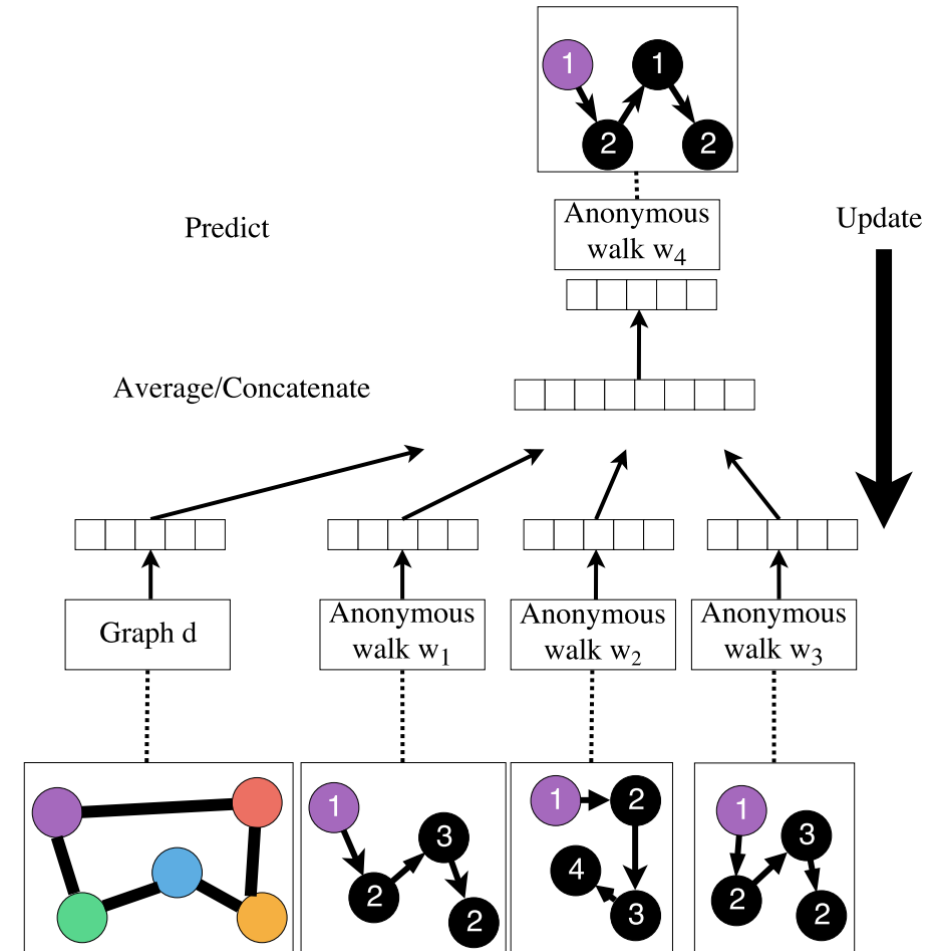
# Learn Walk Embeddings (3)

**Objective:**  $\max \frac{1}{T} \sum_{t=\Delta}^{T-\Delta} \log P(w_t | \{w_{t-\Delta}, \dots, w_{t+\Delta}, \mathbf{z}_G\})$

- $P(w_t | \{w_{t-\Delta}, \dots, w_{t+\Delta}, \mathbf{z}_G\}) = \frac{\exp(y(w_t))}{\sum_{i=1}^{\eta} \exp(y(w_i))}$   All possible walks  
( $\eta$  be number of all possible walk embeddings)
- $y(w_t) = b + U \cdot \left( \text{cat}(\frac{1}{2\Delta} \sum_{i=-\Delta, i \neq 0}^{\Delta} \mathbf{z}_i, \mathbf{z}_G) \right)$ 
  - $\text{cat}(\frac{1}{2\Delta} \sum_{i=-\Delta}^{\Delta} \mathbf{z}_i, \mathbf{z}_G)$  means an average of anonymous walk embeddings in window, concatenated with the graph embedding  $\mathbf{z}_G$
  - $b \in \mathbb{R}$ ,  $U \in \mathbb{R}^D$  are learnable parameters. This represents a linear layer.
  - $\mathbf{z}_i, \mathbf{z}_G$  are learnable.

# Learn Walk Embeddings (4)

- We obtain the graph embedding  $\mathbf{z}_G$  (learnable parameter) after optimization
- Use  $\mathbf{z}_G$  to make predictions (e.g. graph classification)



Overall Architecture

# Summary of Part 3

**We discussed 3 ideas to graph embeddings**

- **Approach 1:** sum/mean/max/hierarchical pooling
- **Approach 2:** Create super-node that spans the (sub) graph and then embed that node
- **Approach 3: Anonymous Walk Embeddings**
  - Idea 1: Sample the anon. walks and represent the graph as fraction of times each anon walk occurs
  - Idea 2: Jointly learn anonymous walks' embeddings and graph embedding

# Today's Summary

We discussed **graph representation learning**, a way to learn **node and graph embeddings** for downstream tasks, **without feature engineering**.

- **Encoder-decoder framework and “shallow” encoding:**
  - Encoder: embedding lookup
  - Decoder: predict score based on embedding to match node similarity
- **Node similarity measure: (biased) random walk**
  - Examples: DeepWalk, Node2Vec
- **Extension to Graph embedding: Node embedding aggregation and Anonymous Walk Embeddings**